

Defense 3: a predetermined acknowledgement delay for DNP3, built in the data plane of a programmable switch

Validated on an Intel Tofino-1 against a physical SEL-751 protection relay

30 July 2026 · 480 physical transactions · `research/case-a-defense3-fixed-ack-delay`

CORRECTED 30 July 2026 after an external audit, then REPAIRED. Every correction the audit demanded that could be verified is applied here and marked [AUDIT]. It found three defects; **all three are now repaired and each is validated on silicon** — R1 (a RESPONSE marking before validation) across 1920 live transactions and Gate 4 case F; R2 (fail-open not generation-qualified) at two budgets, the path now crediting all 64 tokens instead of 1; R3 (a host-injected `0x88C1` entering the queue) via an in-switch forged-frame injector that R3 drops (§7.5–§7.8). The stale-response isolation PASS was withdrawn and **re-established on the repaired build with master-side capture** (§9.8); and the defended-transaction count, the $D = 16$ ms distribution, the fail-open margin, the trap classification and the strength of the headline claims are all corrected. **Every measurement in §10–§11 predates the repairs. All audit remediation items are complete** — including the parser start-state initialization (§7.10, zero warnings across all four builds) and the TCP-sequence-zero repair (§7.9) — and the canonical build is validated on silicon. **Several broader evaluation questions remain open** (not audit-remediation items, not release gates): external-wire adversarial injection (topology-blocked), the cross-transaction generation-wrap reproduction, a hardware-timestamped observer capture, the ACK-retirement boundary sweep, K-minimization, and a full physical core-vs-telemetry parity campaign — see the final-status matrix at the end of §12. Verification: `AUDIT_RESPONSE.md`; repair work: `design/REPAIR_R1_R2_R3.md`, `evidence/{repaired,failopen,inject}/`.

How to read this document. It starts with the physical world — a relay in a substation and someone watching the wire — and narrows, step by step, to nanosecond measurements inside a switch chip. Nothing is assumed: every term is defined where it first matters, every formula is derived, and every number is either measured in this work or marked as inherited. The verdict is stated early and then earned: *the timing a predetermined ACK delay larger than the native CLRT compresses that CLRT distribution by a factor of about 238 on the one relay tested — and the defense that does it is trivially visible to the same observer*. Both halves of that sentence are load-bearing. [AUDIT] An earlier version said the fingerprint was *genuinely destroyed*; with one device, no confusion set and no cross-device classifier that was not supportable.

Final system at a glance the one-screen summary of the finished design, before the investigation.

item	final value
Protected exchange	READ → pure ACK → RESPONSE (a separate-ACK device, the SEL-751)
Transformation	the pure TCP ACK is held until $t_{\text{ACK}} + D$, released independently of the RESPONSE
Ordering	the ACK and an in-window RESPONSE share the Q_{HOLD} FIFO, so ACK-before-RESPONSE is structural
Blocking	$K = 64$ self-recirculating tokens in Q_{BLOCK} , strict-priority over Q_{HOLD}
Fail-open	a generation-labelled <code>reg_failopen</code> note (R2); the budget is the only terminator when no ACK arrives
Final repairs	R1, R2, R3, E1, the TCP-sequence-zero fix (§7.9), and parser initialization (§7.10)
Canonical source	<code>p4/case_a_defense3.p4</code> — R1/R2/R3 unconditional (a no-flag build is the safe program)
Core resources	10/12 ingress stages, 0 egress, critical path 10 (bf-p4c 9.13.2)
Telemetry resources	11/12 ingress stages, 0 egress, critical path 10
Physical evidence	2400 total transactions across three campaigns, 2000 defended
Main result	CLRT standard deviation 2.854 ms → 0.012 ms at $D = 16$ ms ($\approx 238\times$, not flattened to a constant)
Main limitation	the defense is itself detectable in the measured sessions; device anonymity is not demonstrated

item	final value
Release status	canonical artifact frozen and hardware-validated; all audit items closed; optional experiments remain (§12)

Contents

1	The setting, in plain terms	4
2	The vocabulary, defined once	4
3	The leak: why CLRT identifies a device	5
4	Three possible defenses, and the arithmetic that selects one	6
5	How to delay a packet inside a chip that has no timers	6
6	The arithmetic	8
6.1	The deadline, and why D is quantized	8
6.2	The release bias: why the hold is longer than D	8
6.3	The fail-open horizon	9
6.4	The trigger chain	10
7	The implementation, and four hardware traps	10
7.1	Trap one — a large constant in stateful hardware silently does nothing	11
7.2	Trap two — a sign test that is always false	11
7.3	Trap three — four operations per register, and it is a hard error	12
7.4	Trap four — the chip has two pipelines, and a timer fires in both	12
7.5	[AUDIT] Three defects — all three repaired and validated on silicon	12
7.6	[AUDIT] What the repairs cost, and why one of them is refuted	13
7.7	[AUDIT] R2 on silicon, and a fingerprint already in the evidence	14
7.8	[AUDIT] Injecting the adversarial frames, and a defect narrower than it read	15
7.9	[AUDIT] Full-width TCP sequence tracking (the sequence-zero repair)	16
7.10	[AUDIT] Explicit parser metadata initialization (the uninitialized-meta warning)	16
8	The state machine, and the bug that took two attempts	17
8.1	What the state has to express	17
8.2	The bug: a missing exit	17
8.3	The repair that could not be built	17
8.4	The repair that was built	17
8.5	A defect that the repair itself introduced	18
8.6	The final state-transition table	19
9	Validation on synthetic traffic: gates 1 to 4	19
9.1	Gate 1 — does it load, and is the hardware configured	20
9.2	Gate 2 — one transaction, seventeen requirements	20
9.3	Gate 3 — five consecutive transactions, no reset between them	20
9.4	Gate 4A — a response arriving just before the deadline	20

9.5	Gate 4B — a response arriving after the acknowledgement has gone	20
9.6	Gate 4C — the relay never answers	21
9.7	A duplicate response, and an invariant that was being violated	21
9.8	A stale response arriving during a live transaction	21
9.9	Resource cost of the whole thing	22
10	Validation on the real relay	23
10.1	The physical setup, read from the hardware	23
10.2	Stage 2 — the connection only, no request	23
10.3	Stage 3, first attempt — a real failure worth recording	23
10.4	Stage 3, second attempt — the hold works	24
10.5	[AUDIT] The repaired build against the same relay	24
11	The <i>D</i>-sweep campaign, the data, and the analysis	25
11.1	How it was designed, and why	25
11.2	The result	26
11.3	The mechanism held up over the 400 defended transactions	27
11.4	The second analysis, which changes how the first should be read	28
11.5	What survives, precisely	30
12	What may and may not be claimed	30
12.1	Established	30
12.2	Not established, and why	31
12.3	[AUDIT] What the implementation requires, which the earlier text did not disclose	32
12.4	Open work, and what each item blocks	32
12.5	[AUDIT] Final status — what is closed and what remains open	34
12.6	Stated head-on rather than buried	34
13	How to reproduce everything	35
13.1	Software only, no hardware	35
13.2	Compiling	35
13.3	The figures	35
13.4	On hardware	36
14	Every mistake made, and what it cost	36
15	Summary	38

The relay is untouched. No packet byte is modified. Only the TIMING of the ACK leaving port 9 is changed. Port 9 runs at 25 Gb/s, port 64 at 1 Gb/s.




Figure 1: Where the switch sits. The relay and the master are both untouched software; the switch between them runs the defense. Measurement was taken with a host PCAP at the **master interface**, which is a *proxy* for a port-9 wire observer, not the wire observation point itself — the timing numbers are therefore an approximation of what an attacker on the wire would obtain, at $\sim 1 \mu\text{s}$ capture resolution (§12, open-work #14). Port numbers, front-panel positions and link speeds were read out of the live switch configuration, not taken from a configuration file — see §10.

1 The setting, in plain terms

An electrical substation contains **protection relays**: devices that watch the current and voltage on a power line and open a breaker when something goes wrong. A control centre talks to them over a network using an industrial protocol called **DNP3**. A typical exchange is monotonous. The control centre asks “what is your status?”; the relay answers with a list of measurements. This repeats every second or so, indefinitely.

Someone who can see that traffic — a *passive* observer, who cannot inject or alter anything — would like to know **what equipment is installed**. Knowing that a particular cabinet holds an SEL-751 rather than some other relay tells them which vulnerabilities to try, which firmware to study, and what the substation protects. In security terms this is **reconnaissance**: the step before an attack.

Encryption does not solve it, for two reasons. First, a great deal of installed DNP3 is unencrypted. Second and more fundamentally, **encryption hides what messages say, not when they are sent**. If a device always takes about three milliseconds to think before answering, that three milliseconds is visible whether or not the answer is encrypted. Timing is a side channel, and a side channel is not closed by a cipher.

This work closes one specific timing side channel, in the network, without touching the relay and without modifying a single byte of any packet.

2 The vocabulary, defined once

DNP3

The application protocol. For our purposes it runs over TCP, and a *poll* is one request/response pair.

A Class-0 READ

The commonest DNP3 request: “send me your current static data”. It is *read-only*. Nothing in this project ever sent a command that could change the relay’s state; see §12.

TCP acknowledgement (ACK)

TCP is the transport underneath DNP3. When a machine receives data, it sends back a tiny empty packet meaning “got it”. That is an ACK. It is produced by the operating system’s network stack, almost immediately, and carries no application data.

Separate-ACK versus combined-ACK

When a relay receives a READ, two things must happen: TCP must acknowledge the bytes, and the application must compute an answer. Some devices send the ACK at once and the answer later — *two packets*, a **separate-ACK** device. Others wait and put both in one packet — **combined-ACK**. This distinction is the whole ballgame.

Case A / Case B

Case A is a separate-ACK device. The SEL-751 is one. Two packets means a measurable gap between them, and that gap is the target of this work. **Case B** is a combined-ACK device: one packet, no gap, nothing to hide. Case B is out of scope.

CLRT — Cross-Layer Response Time

The gap, on a separate-ACK device, between the transport acknowledgement and the application’s answer. The name says what it is: an interval measured *across* two layers.

Tofino

A programmable switch chip. You write a program stating what the switch does to each packet; the language is **P4**. A Tofino is not a computer. It has no loops, no waiting, and a fixed budget of processing *stages* through which every packet marches in lockstep. Everything in §5 follows from those constraints.

D The one parameter of this defense: how long to hold the ACK back.

Formally, for one transaction, with t_{READ} , t_{ACK} and t_{RESP} the instants at which the observer sees the three packets:

$$\text{CLRT} = t_{\text{RESP}} - t_{\text{ACK}}, \quad a = t_{\text{ACK}} - t_{\text{READ}}, \quad \text{total} = a + \text{CLRT}. \quad (1)$$

Throughout, a is the relay’s own acknowledgement latency and $c = \text{CLRT}$ is the quantity we are trying to hide. Keeping these three symbols straight is the whole of §11: a defense can shrink c and yet leave the information intact somewhere else in (1).

3 The leak: why CLRT identifies a device

When a relay receives a READ, its TCP stack acknowledges the bytes within a fraction of a millisecond: cheap, mechanical work done by the network stack. Then the application firmware wakes up, gathers measurements, formats a DNP3 response and sends it. How long that second step takes depends on the relay’s processor, its firmware, its scheduler and how much data it must format — in other words, **on what the device is**. Different models differ. That is what makes CLRT a fingerprint.

The prior work that established this (Formby et al.) treats CLRT as a *physical property of the device* and uses it to identify equipment on a network. So the defensive objective this project inherited is exactly: **make the observed CLRT stop revealing the device**. That is the objective this report grades the defense against; grading it against a broader one (“make the device unidentifiable”) is a different and much harder question, treated honestly in §12.

Measured on the real SEL-751 with no defense running, 80 polls:

quantity	value
CLRT median	2.828 ms
CLRT standard deviation	2.854 ms
CLRT 95th percentile	12.222 ms
CLRT maximum	13.175 ms

The spread is the information. A device whose CLRT is usually about 2.8 ms with excursions to 13 ms looks different from one that is always at 0.4 ms. The *standard deviation* is the number to follow through this whole report, because a defense that leaves the spread intact has hidden nothing, however much it moves the median.

4 Three possible defenses, and the arithmetic that selects one

Given a two-packet exchange and a switch in the middle, there are exactly three things the switch can do to the gap.

Defense 1 — hold the ACK until the response arrives, then release both.

The gap collapses to nearly zero. Already built in this project. Its flaw is structural: the switch releases the ACK *because* the response arrived, so the ACK’s timing now depends on the response. The interval $t_{\text{READ}} \rightarrow t_{\text{ACK}}$ becomes $a + c$, and c is still there — moved, not removed.

Defense 2 — forward the ACK at once, hold the RESPONSE to a fixed deadline G .

Also already built, and proven on this hardware against this relay. The observed CLRT becomes G , a constant. Its cost is latency: the answer is genuinely late by up to G .

Defense 3 — this work.

Forward nothing on demand. Hold the ACK to a fixed, *predetermined* instant $t_{\text{ACK}} + D$, chosen before the response is known and released independently of it. The response, if it arrives during the hold, queues *behind* the ACK and leaves just after it.

The case for Defense 3 is arithmetic, not intuition. Writing Δ_{release} for the residual ACK-to-RESPONSE gap left by the release machinery (this is the “ δ ” used loosely below — and it is **not a universal constant**: it is a capture-dependent distribution folding FIFO release, egress scheduling, serialization, NIC handling and host timestamping, with a median of about $32 \mu\text{s}$ in this campaign, not an architectural floor):

defense	observer sees READ→ACK	observer sees CLRT	is c visible?
none	a	c	yes, directly
Defense 1	$a + c$	≈ 0	yes — inside READ→ACK
Defense 2	a	G	no
Defense 3	$a + D$	$\max(c - D, \delta)$	no, when $D > c$

The last row is the point. Because the release instant is fixed *in advance* and does not depend on the response, c appears in **neither** observable, provided D exceeds c :

$$\text{CLRT}_{\text{out}} = \begin{cases} c - D + \epsilon_{\text{direct}}, & c > D \quad (\text{response forwarded directly, 0 passes}) \\ \Delta_{\text{release}}, & c \leq D \quad (\text{response held behind the ACK}), \end{cases} \quad (2)$$

$$(t_{\text{READ}} \rightarrow t_{\text{ACK}})_{\text{out}} = a + D_{\text{realized}} + \epsilon_{\text{release}}.$$

The conceptual approximation $\text{CLRT}_{\text{out}} \approx \max(c - D, \Delta_{\text{release}})$ is useful, but the exact statement is the piecewise form above: Δ_{release} is the *measured* residual distribution (median $\approx 32 \mu\text{s}$ here), and READ→ACK is $a + D_{\text{realized}} + \epsilon_{\text{release}}$, not exactly $a + D$.

Consequence that drove the whole evaluation. The right value of D is not “the average CLRT”. Equation (2) conceals a transaction only when $D > c$, so **only transactions faster than D are hidden**. A D set at the median hides about half of them. This is why §11 sweeps D instead of picking one.

Figure 2 draws all four situations on one axis, with the measured medians ($a = 0.45 \text{ ms}$, $c = 2.85 \text{ ms}$) so the geometry is to scale.

5 How to delay a packet inside a chip that has no timers

This is the hardest idea in the project and every later number depends on it.

A Tofino processes a packet through a fixed pipeline and sends it out. **There is no instruction meaning “hold this packet for two milliseconds”**. There is no timer you can attach to a packet, no sleep, no programmable delay queue with a deadline. Packets go in and come out.




Figure 2: The same exchange with no defense and under each of the three defenses, to scale with the measured medians. The orange span in each row is the interval that carries the secret. Under Defense 1 the secret has simply moved from the CLRT into READ→ACK. Under Defense 3 neither interval contains c , provided $D > c$. ACK and RESPONSE are drawn in two lanes per row because under Defenses 1 and 3 they leave within microseconds of one another and would otherwise coincide. Source: `figures/src/fig4_timelines.py`.

What the chip does have is a **traffic manager** with queues and **strict priority**. If queue 7 is strict-priority over queue 1, then for as long as queue 7 holds anything at all, queue 1 is not served. That is the entire lever, and the mechanism is built from it:

1. Put the ACK to be delayed into Q_{HOLD} (queue 1, low priority) on an internal loopback port.
2. Fill Q_{BLOCK} (queue 7, strict high priority) with dummy packets — **blocker tokens** — so Q_{HOLD} is starved and the ACK cannot leave.
3. Have each token, when it is served, read the clock, decide whether the deadline has passed, and if not **send itself around the loop again**. A recirculating token keeps Q_{BLOCK} non-empty.
4. Once the deadline passes, every token stops recirculating and is dropped. Q_{BLOCK} empties, Q_{HOLD} is served, and the ACK leaves.

The delay is therefore produced by a **self-sustaining crowd of dummy packets that gets out of the way at the right moment**. Nothing waits; everything is a packet being processed normally, over and over.

Three consequences matter.

Why 64 tokens and not one? A single token leaves gaps. After it is served it takes a fraction of a microsecond to come back around, and in that window Q_{BLOCK} is empty and the ACK escapes. Enough tokens must be in flight that the queue is *never* momentarily empty. $K = 64$ was established as sufficient by earlier work in this project, and every measurement here confirms it. It is *not* claimed to be minimal.

Where do the 64 tokens come from? The switch makes them itself. Tofino has a **packet generator** that can be triggered by a pattern in a recirculated packet. When a READ arrives and arms a transaction, the switch mirrors a small tagged copy of it to the generator's port; the generator recognises the tag and emits 64 tokens. No external host is involved — which matters, because a defense that needs a server feeding it packets is not a defense.




Figure 3: (a) The construction. Q_{BLOCK} sits at strict priority 7 and Q_{HOLD} at priority 0 on the same internal loopback port, so Q_{HOLD} is never served while a single token remains. (b) What happens *after* the deadline, at nanosecond scale, measured on the physical relay. The hold is D plus these three terms, and D itself is 1 999 691 ns — three orders of magnitude larger, which is why it cannot be drawn on the same axis. Source: figures/src/fig2_mechanism.py.

What if something goes wrong? Each token carries a **budget**: a maximum number of loops. If the deadline somehow never arrives, tokens exhaust their budget and stop anyway. This is the **fail-open** path — the defense gives up and lets the packet through rather than holding it forever and breaking the TCP connection. §6 gives the arithmetic that sets the budget.

6 The arithmetic

Four numbers govern the mechanism. All are measured or derived; none are guessed.

6.1 The deadline, and why D is quantized

The switch stores the deadline as a 32-bit word. Its low byte is reserved as an armed/not-armed flag, so the time itself lives in the upper 24 bits, counted in units of 256 nanoseconds — one *tick*. To hold for D milliseconds:

$$\text{ticks} = \left\lfloor \frac{D \times 10^6}{256} \right\rfloor \quad (\text{round DOWN, so } D_{\text{realized}} \leq D_{\text{requested}}), \quad \text{word} = \text{ticks} \ll 8. \quad (3)$$

The left shift puts a zero in the low byte so that the armed flag survives the later addition. For $D = 2$ ms:

$$\text{ticks} = \lfloor 2\,000\,000/256 \rfloor = 7\,812, \quad \text{word} = 7\,812 \ll 8 = 1\,999\,872 \text{ ns},$$

so $D_{\text{realized}} = 1.999872$ ms and the shortfall is 128 ns. **Every “ $D = 2$ ms” in this report means 1 999 872 ns.** That is quantization, not a defect. The admissible range of D is **no longer a fixed clamp**: the parameter policy computes D_{max} from the fail-open horizon (below), and the original 40 ms clamp (inconsistent with $H = 30.802$ ms) has been removed. poll.

6.2 The release bias: why the hold is longer than D

When the deadline passes, the tokens do not vanish together. Each one only learns of it the next time it is *served*. Draining K tokens out of the queue at the loop’s packet rate takes

$$\tau = \frac{K}{r_{\text{dp8}}} = \frac{64}{37.4 \times 10^6 \text{ pps}} = 1\,711 \text{ ns}, \quad (4)$$

where 37.4 million packets per second is the *measured* rate of the 25 Gbit/s internal loopback for the small frames involved. So the ACK is released not at the deadline but at approximately deadline $+ \tau$.

This matters for honest measurement. Score the hold against D and you see a systematic $+1\,711$ ns error on *every* transaction, and you will be tempted to call it jitter. The correct target is $D + \tau$, and every deadline number in this report is given both ways: raw against D , and corrected against $D + \tau$.

And τ was verified independently, which is the part that matters. Originally τ was a model whose only evidence was the residual it removed — circular reasoning. Instrumentation was added to timestamp the *first* and *last* token terminations, so the drain could be measured directly:

$$\text{predicted } \tau = 1\,711 \text{ ns} \quad \text{measured drain} = 1\,692\text{--}1\,696 \text{ ns} \quad \text{agreement } 1.1 \text{ \%}.$$

The full decomposition of one hold, from the synthetic gate, with each term measured separately and nothing inferred:

$$\underbrace{2\,001\,505}_{\text{hold}} = \underbrace{1\,999\,763}_{D, \text{ tick-quantized}} + \underbrace{1\,692}_{\text{drain}} + \underbrace{27}_{\text{release tail}} + \underbrace{23}_{\text{detection}} \text{ ns.} \quad (5)$$

[AUDIT] Two different quantities were both called “the release tail”. They are not the same event and they differ by three orders of magnitude. From here on: *internal release tail* = 26–27 ns, the last token termination to the held ACK’s loopback return, inside the switch; *external ACK*→*RESPONSE floor* $\approx 32 \mu\text{s}$, the gap an observer captures at the master. The second is not the first scaled up — it also contains switch output queuing, frame serialization, link traversal, NIC processing and host capture behaviour. Every later “32 μs ” means the external floor.

[AUDIT] The drain model is off by one. The interval from the *first* termination to the *last* spans $K - 1$ gaps, not K :

$$\frac{K-1}{r} = \frac{63}{37.4 \times 10^6} = 1\,684.5 \text{ ns}, \quad \frac{K}{r} = 1\,711.2 \text{ ns}, \quad \text{measured } 1\,692\text{--}1\,696 \text{ ns}.$$

The measurement fits $(K-1)/r$ better (9.5 ns versus 17.2 ns). K/r remains the right figure for the reservoir’s full circulation period and the release bias it predicts is still correct to about 1 %, but the claim that the measurement “independently verifies K/r ” was imprecise.

6.3 The fail-open horizon

Each token may loop at most B times. With K tokens sharing the loop, the wall-clock time before the last one gives up is

$$H = \frac{BK}{r_{\text{dp8}}} = B\tau = 18\,000 \times 1\,711 \text{ ns} = 30.802 \text{ ms.} \quad (6)$$

[AUDIT] This constraint was written against the wrong quantity. The budget starts when the reservoir is created, a few hundred nanoseconds after the READ — but the deadline is $t_{\text{ACK}} + D$, and t_{ACK} can be milliseconds later. The horizon must clear the relay’s own acknowledgement latency too:

$$H > a + D + t_{\text{detection}} + t_{\text{drain}} + t_{\text{tail}}.$$

Measured against the campaign’s own worst-case a per arm:

arm	worst observed a	$a + D$	margin $H/(a + D)$
$D = 1$	4.608 ms	5.608 ms	$5.49\times$
$D = 2$	3.399 ms	5.399 ms	$5.71\times$
$D = 4$	5.651 ms	9.651 ms	$3.19\times$
$D = 8$	1.509 ms	9.509 ms	$3.24\times$
$D = 16$	4.673 ms	20.673 ms	1.49$\times$

An earlier version quoted $8.8\times$, computed as $H/3$ ms from a stale design point with a omitted entirely. **The true worst-case margin in this campaign was $1.49\times$** , at $D = 16$ ms. Fail-open still never fired, but the headroom is a third of what was claimed, and a transaction with the historically observed 22 ms READ→ACK would have exceeded the budget at $D = 16$.

The original parameter range was arithmetically impossible — and has been fixed. The *original* control plane clamped D at a fixed 40 ms while $B = 18\,000$ gives $H = 30.802$ ms; at that clamp the budget expires *before* the deadline can arrive. **The final parameter policy removes the fixed clamp** and computes $D_{\max} = H - a_{\text{bound}} - t_{\text{detect}} - t_{\text{drain}} - t_{\text{tail}} - M \approx 24.8$ ms (`control/parameter_policy.py`), so $D = 16$ ms is admitted and $D = 40$ ms is **refused**, verified on silicon. This resolves audit open-work #5. For reference, the horizon window:

$$\frac{H}{20.673} = 1.49 \times \text{ (thin but clear), } \quad \frac{200 \text{ ms}}{H} = 6.5 \times, \quad \frac{H}{40 \text{ ms}} = 0.77 \times \text{ **INFEASIBLE.**}$$

An inherited comment in the code assumed $10\,\mu\text{s}$ per loop, giving a horizon $5.8\times$ wrong. It was replaced by (6), because a wrong model gives a wrong answer the moment K , B or the port speed changes. **Fail-open fired zero times during the healthy, deadline-governed gates and the physical D-sweep campaigns** — the outcome you want: the safety net exists and was never needed there. It *was* deliberately exercised in the READ-only fail-open and K-sweep experiments used to validate R2 (§7.7), where a READ with no ACK arms no deadline and the budget is the only terminator by design.

6.4 The trigger chain

How long after the READ does the reservoir actually exist? Measured over **100 clean trials**, total spread **4 ns**:

step	time
READ arrives → tagged copy returns to the generator	688 ns
generator recognises the pattern → first token admitted	11 ns
all 64 tokens admitted (the burst itself)	516 ns, ≈ 8 ns per token
READ → full 64-token reservoir	1 215 ns

Compare that with the relay’s own acknowledgement latency — the earliest moment there is anything to hold — measured at 0.453 ms median, 0.400 ms minimum. The reservoir is ready roughly **330 times sooner than it needs to be** (Figure 4). There is no race, and the cold first trial after a program load sits inside the warm distribution, so there is no warm-up cost either.

7 The implementation, and four hardware traps

The switch program is `p4/case_a_defense3.p4`. Its shape is simple: decide in the parser what each packet *is*, resolve every remaining condition in *one* table lookup, then act. What is not simple is the hardware. Four separate traps were found.

[AUDIT] They are not all the same kind of thing. An earlier version said the compiler “accepted all four without complaint”, which contradicts §7.3, where it emits a hard error. Classified honestly:

#	trap	what the toolchain did
7.1	large constant in a stateful ALU	confirmed silent target anomaly — compiled, ran, wrote nothing, no diagnostic
7.2	unsigned $v < 0$	programmer type error with a missing diagnostic. $v < 8w0$ on a $\text{bit}<8>$ is <i>correctly</i> false; the compiler is not wrong, it simply never warned that the predicate is vacuous
7.3	a fifth RegisterAction	hard compiler error — loud and immediate
7.4	a timer firing in both pipes	documented target behaviour we had not accounted for

Only 7.1 is a case of the hardware doing something other than what the program said.




Figure 4: The trigger chain against the deadline it must beat, on a logarithmic axis because the two live three orders of magnitude apart. The reservoir is complete 1 215 ns after the READ; the earliest the relay’s own ACK can arrive is 400 000 ns. The margin is $\approx 330\times$, which is why the arming path never races the traffic it protects. Source: `figures/src/fig6_trigger.py`.

7.1 Trap one — a large constant in stateful hardware silently does nothing

The switch keeps a small amount of memory a packet can read *and* modify as it passes: a **register** with a tiny attached processor, the **stateful ALU** (SALU). One register, `reg_tag`, holds the “which transaction is currently live” marker.

The code said: *if the marker says idle, write my transaction’s identity into it*. Idle was originally encoded as `0xFF`. On hardware, **the write never happened**. The transaction never became live, so all 64 tokens were rejected and the relay’s ACK was refused. But the same tiny processor also *returns* a value, and that path worked perfectly — so a counter reported “a new transaction was armed” while the register said no transaction existed. That contradiction produced two wrong diagnoses before the compiled assembly was read:

```
sub hi, phv_lo, lo      the RETURN value — worked
equ lo, lo, -255        the PREDICATE — never true
alu_a cmplo, lo, phv_lo the conditional write, therefore never executed
```

Changing idle from `0xFF` to `0x00` made the predicate `equ lo, lo` — a comparison against zero, needing no embedded constant — and the write began to commit. Tokens admitted went from **0** to **64**.

A correction against ourselves, which is why a probe exists. The obvious explanation is “255 is too wide for the instruction’s constant field”. `p4/probe_salu_immediate.p4` tests that by comparing thirteen registers against $K \in \{1, 2, 7, 8, 15, 16, 63, 64, 127, 128, 192, 254, 255\}$. **The compiler emits `equ lo, lo, -K` for every one of them, identically, with no error and no warning.** So the assembly *cannot* distinguish a safe constant from an unsafe one, and the width explanation is an inference consistent with the evidence, not a proof. The repair is confirmed behaviourally on silicon; the mechanism is not.

The usable rule is structural, not an inspection: never compare stateful state against a large constant. Compare against zero, or against a value carried in the packet. It is enforced by a test, `analysis/test_tag_domain.py`, not by discipline. After the repair, exactly *one* constant comparison remains anywhere in the program — against 2 — and that one is proven working on hardware.

7.2 Trap two — a sign test that is always false

Later work needed “is the top bit of this byte set?”, which is naturally written as “is this value negative”. Written the obvious way on an unsigned byte, `if (v < 8w0)` compiles fine and emits `lss.u lo, lo`. But `lss.u` is

an *unsigned* less-than-zero: it is never true. The compiler reported success. With an explicit signed cast, `if ((int<8>)v < 8s0)` emits `lss.s lo, lo`, which is correct.

[AUDIT] This is not a miscompile. An unsigned less-than-zero really is always false; P4’s semantics are correct and the fault is in the type that was written. What the toolchain failed to do was *diagnose* a predicate it could prove vacuous. Still, one genuine silent anomaly and one undiagnosed type error in the same small piece of hardware was enough to stop trusting inspection, so `analysis/assert_salu_asm.py` now **fails the build** if the compiled assembly for the load-bearing predicates contains `lss.u` or is missing the expected instructions. It is mutation-checked: reverting the cast makes the compiler exit 0 while the assertion exits 1.

7.3 Trap three — four operations per register, and it is a hard error

Adding a fifth way of touching `reg_tag` produced error: too many RegisterActions attached to the Register; the target architecture limits the number to 4.

This forced a genuinely better design. One of the five was a plain read, used by the blocker tokens — and a read is just a modification that adds zero. So the read and the “mark this transaction” operation were merged into a single operation whose increment arrives in the packet’s metadata: `0x50` to mark, `0` for a plain read. Four operations, nothing lost.

7.4 Trap four — the chip has two pipelines, and a timer fires in both

An early synthetic test emitted three events and observed six. The chip has two pipelines, and enabling a packet-generator application device-wide arms it in *both*. A pattern-triggered application is masked from this — only one pipeline sees the trigger — but a timer-triggered one fires everywhere it is armed. Every generator write is now scoped to one pipeline. From the same session: `pgrep -f bf_switchd` over-counts (it returned 3 for one process); `pgrep -cx` is correct.

7.5 [AUDIT] Three defects — all three repaired and validated on silicon

An external audit found three defects — two state-ordering (defects 1 and 2) and a host-injected-token admission path (defect 3) — and reading the source confirmed all three. Since then **all three have been repaired and each validated on silicon**; defect 2’s obvious one-operation repair was shown *not* to fit and a second-register repair was built instead (§7.6–§7.7).

defect	repair	status
1 — a RESPONSE marks before its identity is checked	R1	REPAIRED. 10/12 live core, 11/12 live+telemetry and synthetic, critical path 10. Validated on silicon in the synthetic build (Gate 2 PASS, Gate 3 PASS 10/10, Gate 4 PASS on all six cases) and run against the physical relay for 960 transactions with the hold, the CLRT compression and the ordering invariant all unchanged (§10.5).
2 — fail-open retirement is not generation-qualified	R2	REPAIRED, VALIDATED ON SILICON, AND RUN ON THE LIVE BUILD (§7.7, §10.5). Fail-open now credits all 64 tokens instead of 1, <code>reg_tag</code> survives, the next transaction still arms (28/28 trials at two budgets), and 960 live transactions against the relay show no harm. Single-generation only; the <i>foreign</i> -token case is model-checked, not produced on hardware.
3 — a host-injected <code>0x88C1</code> frame enters the priority queue	R3	REPAIRED and DEMONSTRATED ON SILICON (§7.8). A forged <code>0x88C1</code> token injected in-switch is dropped at the fresh stage and never reaches the loopback; without R3 the same frame enters. Zero resource cost.

The original physical campaign — §10.1–§10.4 and the §11 D-sweep, every number in those results — was collected on the UNREPAIRED build, with all three defects present. The repaired builds appear only in §10.5 (the two R1+R3 and R1+R2+R3 live campaigns totalling 1920 transactions). The repairs do not retroactively change any earlier measurement; they change what the mechanism does next time.

The rule the two state-ordering defects break: state is written before it is validated. The switch resolves a packet’s conditions across pipeline levels, and in both cases the register write happens at level 2 while the test that authorises it resolves at level 3.

Defect 1 — a RESPONSE marks the transaction before its identity is checked.

level	line	what runs
1	1924–1932	class driver sets meta.tag_val = TAG_PENDING_DELTA
2	1973–1978	tag_read_or_mark.execute(0) — the write happens here
3	1990	tbl_state_decode.apply() — seq / ack / port resolve here

CLASS_RESP is assigned on direction, session and DNP3 framing alone, and the stateful ALU’s own guard is `(int<8>)v < 8s0` — a test on the *stored* state, not on *this packet’s* validity. So a correctly framed response on the tracked session with a **wrong TCP sequence** still marks the live transaction. The legitimate response then reads the pending marker, is treated as a duplicate and is **dropped**, and the acknowledgement’s release declines to retire — leaving the transaction stuck until fail-open.

Defect 2 — a foreign zero-budget token can retire the current transaction. At line 1938 a dequeued blocker with `hdr.ib.seq == 0` sets `meta.tag_val = TAG_INACTIVE`; `tag_rmw` commits it at line 1985, guarded only by `tag_val`; and the generation check is at line 1990. The documented *stale > deadline > budget* priority is evaluated in the action block, one level after the write has committed, and a later table cannot undo a stateful-ALU write. Compounding it, the parser forces EtherType `0x88C1` to `ROLE_BLOCK` from *any* topology port, with a legacy branch that enqueues such a frame into the strict-priority queue — so an injected token with a zero budget reaches this path. The threat model is passive, so that injection is outside the modelled adversary, but a production build should not carry it.

Neither defect was observed firing. Across the 400 defended physical transactions, duplicate suppressions were 0, stale terminations were 0 and fail-open fired 0 times. They were latent, not manifested — but “not observed” is not “cannot happen”, and defect 1 invalidated a test this report previously reported as passing (§9.8).

7.6 [AUDIT] What the repairs cost, and why one of them is refuted

R1 — authorise the marker before writing it. It fits. The RESPONSE rows of `tbl_state_decode` mask `tag_diff` out entirely, so the RESPONSE verdict never depended on `reg_tag` at all: it depends only on the three session-tracker differences, and all three are produced *before* the tag access. The same conjuncts therefore resolve one level earlier and choose the marker delta; an unauthorised RESPONSE carries delta 0, making the identical stateful operation a pure read. `reg_tag` keeps its placement and its four operations. Cost: one table and one dependency level, $9 \rightarrow 10$ ingress stages, critical path $8 \rightarrow 10$. Stage count now equals critical path, so the program is dependency-bound at 10; with the telemetry registers it would sit at 11/12.

R2 — generation-qualify the fail-open retire. Three refuted attempts, then a repair. The two operations look mergeable into one stateful arm — *if idle, arm* and *if it is mine, retire* are the same shape. They are not. `p4/probe_failopen_qualification.p4` reduces the first two walls to a minimal program:

build	result
four operations, unmerged	compiles — the control
merged into one arm	error: The input meta.tag_alt to stateful alu reg_tag is not allocated in a valid region on the input xbar to be a source of an ALU operation
kept separate (a fifth operation)	error: too many RegisterActions attached to the Register

A third attempt, packing both operands into one 16-bit field, named the real constraint outright: `error: Ingress.reg_tag requires more than 2 PHV inputs`. **reg_tag’s stateful ALU has a budget of two**

PHV inputs shared across all four of its operations, and it was already full (`gen_in` and `tag_val`). No third source can be added, however packaged.

The repair works by not needing one. Two observations. The fail-open write never released anything — the held ACK leaves because the budget-zero token *drops itself* and Q_{BLOCK} empties — so its only job was to let the **next** READ arm. And on the ARM path `meta.tag_val` is dead. So the note rides on `tag_val`, an operand `reg_tag` already has, and **the generation qualification moves from the producer to the consumer**: a budget-zero token records the generation *it* carries in `reg_failopen`, and the next READ arms if `reg_tag` is idle *or* equals the noted generation. A foreign token’s note names a generation that is not the live one, so it can never authorise anything. The note is cleared as it is read, so it authorises at most one arm.

The note is an observation, not proof of ownership. A stale or foreign budget-zero token also writes `reg_failopen` (the injection matrix shows a foreign `0xC1` token recording `note = 0xC1`); safety comes entirely from the *consumer*, which arms over the note only if `reg_tag == reg_failopen`. Three invariants must stay tested: every READ consumes or clears the note; a mismatching note cannot authorise; and an old matching note cannot survive until the generation value is reused. So R2 stays safe independently of R3, which closes only the external forged-token route.

Cost: none on top of R1 and R3 — 11/12 stages, critical path 10, identical without it. Verified offline three ways: the compiled assembly is *asserted* to contain both comparisons and a write predicated on their OR (`alu_a (cmplo | cmphi)`), because a predicate that compiles and is never true is exactly the trap of §7.1 and §7.2; the state model gained 321 assertions over all sixteen generations and all ordered foreign pairs; and the suite is mutation-checked (dropping the note comparison gives 16 failures, arming unconditionally 224, making the note reusable 16). **R2 was subsequently loaded and validated on silicon at two budgets; see §7.7.**

R3 — refuse host-injected blocker frames. Free. 9 ingress stages, critical path 8, resources bit-identical to baseline. It matters more than its size: it removes the only *practical* route to defect 2. Without an injectable token, reaching defect 2 requires a blocker to outlive its own generation’s deadline, never observed in 25 600 tokens.

A repair that broke the mechanism, caught by Gate 2 on the first transaction. R1’s first version gave its authorisation table a catch-all default action setting `tag_val = 0`. That reaches every packet class, and for every class other than the RESPONSE the tag arm is `tag_rmw`, whose write is guarded by `tag_val != TAG_NO_WRITE` — so it became an unconditional write of `TAG_INACTIVE`. The READ armed the generation and the mirrored trigger clone, returning ≈ 700 ns later, wiped it: `PKTGEN_ADMIT=0`, `PKTGEN_DROP=64`, `ACK_REJECT=1`. Fixed by making “not authorised” a `CLASS_RESP` entry rather than the table default. The defect had already passed 2 354 offline assertions and a compile-fit check: **the offline model covers the state machine, not which table default reaches which packet class.**

7.7 [AUDIT] R2 on silicon, and a fingerprint already in the evidence

A correction first. I wrote that fail-open “has fired 0 times in every campaign, so the path has never executed on silicon”. True of the gates and both *D*-sweeps, where the deadline always beat the budget — but `--check2` is READ-only by construction, so no ACK arrives, no deadline is armed, and the tokens can *only* terminate on the budget. The path had been executing all along; what had not been done was reading what it recorded.

The defect was visible in evidence collected a day earlier. Unrepaired build, 60 trials, every one: `BLOCK_TERM_TMO = 1`, `BLOCK_TERM_STALE = 63`, `reg_tag` afterwards = 0. One token terminates on the budget and sixty-three are credited as *stale*. That 1/63 split **is** the defect — the first token to reach budget zero writes `TAG_INACTIVE` with no generation test, and the other 63 then compute `gen_in - 0  $\neq$  0`, read as foreign and are dropped. Both outcomes are “token terminated”, so nothing looked wrong.

	unrepaired	R2, $B = 18\,000$	R2, $B = 500$
<code>BLOCK_TERM_TMO</code>	1	64	64
<code>BLOCK_TERM_STALE</code>	63	0	0
<code>BLOCK_LOOP</code>	1 152 000	1 152 000	32 000
<code>reg_tag</code> afterwards	0 (cleared)	0xC0 (preserved)	0xC0 (preserved)
next trial <code>ARM_FRESH / PKTGEN_ADMIT</code>	1 / 64	1 / 64	1 / 64

28 trials per arm, every prediction holding in all of them. The recovery property survives, which was the risk in not clearing `reg_tag`: the next READ arms through the note and its reservoir is admitted in full, `ARM_BUSY` = 0 throughout. And the budget arithmetic is confirmed exactly — `BLOCK_LOOP` = $K \times B$ on the nose at both budgets, which is what $H = BK/r$ rests on.

One guard had to be scoped. The shrunk budget was first *refused*: $H = 0.856 \text{ ms} \leq a_{\text{worst}} + D = 24 \text{ ms}$, i.e. the budget would fire during a legitimate hold. Correct in general — that is the §6.3 failure mode — but the wrong test for a READ-only trial, where no ACK arrives and there is no hold to cut short. It is now gated behind an explicit `--read-only-trial`; the general case is untouched. **A safety check that fires on the one scenario a mechanism exists for is usually mis-scoped, not too strict, and the fix is to narrow its precondition rather than remove it.**

The single-token case is pinned down by §7.8’s K -sweep: on the unrepaired build a reservoir of *one* native token gives `TMO` = 1, `STALE` = 0, `reg_tag` cleared, so a single budget-zero token does write the tag and $1/(K-1)$ is the mechanical cascade. R2 turns every K into K budget terminations, 0 stale, tag preserved.

These trials are single-generation. Every token carries the live generation, so they exercise note-and-recover and the within-transaction accounting, **not** the *cross*-transaction case — a token from a retired transaction clearing a *different* live one — which needs the generation-wrap coincidence and remains model-checked.

7.8 [AUDIT] Injecting the adversarial frames, and a defect narrower than it read

The three cases the relay will not produce — a mis-sequenced response (R1), a foreign token at budget zero (R2), an injected `0x88C1` frame (R3) — all reduce to putting a chosen frame on a switch port, which the lab cannot do from a host (no raw socket on the master, no host on the relay leg). The injector is therefore built *inside the switch*, under a flag `D3_INJECT`: a parser path that treats a tagged generator frame as a fresh host-injected `0x88C1` token carrying an attacker-chosen generation and budget. It compiles at the same 11/12 and is a no-op when the flag is absent.

R3 is now demonstrated on silicon — it had been completely unexercised. Under R3 the forged token is dropped at the fresh stage and never reaches the strict-priority queue; without R3 it enters (`BLOCK_TERM` on the loopback). **R2’s note mechanism is shown executing:** with R2 the injected token’s generation is recorded in `reg_failopen` and `reg_tag` is preserved.

[AUDIT] A counter that misreported the drop, corrected and re-verified. The first injector build counted the R3 drop with `BLOCK_ENQ`, the counter that elsewhere means *residence in Q_BLOCK* — so a dropped frame wrongly incremented an “enqueued” tally. The P4 now increments a distinct `BLOCK_REJECT` on the R3 drop, and `BLOCK_ENQ` fires only on an accepted `to_block()`. Re-run on silicon (foreign gen `0xC1`, seq 0, `0xC0` live): the R1+R2 build gives `{BLOCK_ENQ: 1, BLOCK_TERM_STALE: 1}` with `reg_failopen` = `0xC1` (the accepted token is enqueued, then stale-dropped at dequeue, and R2 noted its generation), while R1+R2+R3 gives `{BLOCK_REJECT: 1}` alone — no enqueue, no dequeue-side termination, `reg_failopen` = 0. The drop behaviour is unchanged; only the accounting is now correct.

A puzzle the injector raised, and a microbenchmark that resolved it. The *injected* token above did not clobber `reg_tag` on any build — which seemed to say the write never fires. That was an **injection-harness artifact**. A fail-open K -sweep on the **native** reservoir (READ-only, so the budget is the only terminator) settled it:

K	1	2	4	8	16	32	64
budget-expiry (TMO)	1	1	1	1	1	1	1
stale	0	1	3	7	15	31	63
<code>reg_tag</code> after	0	0	0	0	0	0	0

TMO = 1, STALE = $K - 1$ at every K , and `reg_tag` is cleared even at $K = 1$. The write the audit predicted *does* fire, from a single **native** token, and the $1/(K-1)$ cascade is its mechanical consequence: the first budget-zero token clears the tag, and every later token then reads foreign and terminates stale. The injected token was the anomaly (a frame forged through the legacy path with `seq` = 0 from the start does not traverse the same write as a token that was admitted and looped its budget down). **The defect is real and reproduces at $K = 1$;** the earlier “a single token does not clobber” was wrong about the native path.




Figure 5: The fail-open reconciliation. **(a)** On the unrepaired build the budget-zero terminations are 1 (budget) and $K - 1$ (stale) at every reservoir size, including $K = 1$ — one native token clears `reg_tag` and the rest read foreign, so the defect is not an emergent effect of size. **(b)** R2 turns the same event into K budget terminations and 0 stale, and preserves `reg_tag`. Source: figures/src/fig9_ksweep.py.

It is a *within*-transaction effect — every token carries the live generation, so clearing `reg_tag` is that transaction’s own fail-open, and the harm is the corrupted accounting ($1/K - 1$ instead of $K/0$) and lost reservoir ownership. R2 fixes it (§7.7). The *cross*-transaction clobber — a token from a retired transaction clearing a *different* live one — is a separate claim that still needs the generation-wrap coincidence and remains model-checked; but the write it depends on is now confirmed real and single-token on silicon.

7.9 [AUDIT] Full-width TCP sequence tracking (the sequence-zero repair)

To match a response to the request that armed it, the switch tracks the relay’s expected TCP sequence position in a one-entry register `reg_exp_relay_seq`. The *earlier* design updated it with a *value* sentinel, `if (meta.seq_w != 32w0) { store }`. That reads plausibly, but **TCP sequence 0 is a valid position** — it recurs each time the 32-bit sequence space wraps — so on the wrap the tracker silently declined to update, kept a stale expectation, and would forward the next relay frame *unprotected*. Rare, but structurally wrong.

The repair splits the single read-modify-write into two class-selected actions: a **writer** `exp_seq_w` that **stores unconditionally** (so sequence 0 lands), and a read-only **reader** `exp_seq_r`. Which runs is chosen by the packet class (`meta.sess == SESS_MASTER`) at an MAU gateway, not by the value. It costs **nothing**: the register keeps its two stateful-ALU PHV inputs (`hdr.tcp.seq_no`, `meta.seq_w`), the pattern is a verbatim clone of the `reg_exp_ack` writer/reader split already on silicon, and the compiled footprint is **resource-neutral** (core still 10/12, path 10; the one added logical table is absorbed in the existing tracker stage). The assembly confirms it: `exp_seq_w` emits `sub hi, phv_lo, lo ; alu_a lo, phv_hi ; output alu_hi` — an unconditional store with no value predicate. Gate 2 passes unchanged, so the normal (`seq ≠ 0`) path is untouched; the only behavioural change is that a post-wrap sequence-0 relay frame now qualifies. Evidence: `evidence/final_silicon/*/seqzero_B_validation.md`.

7.10 [AUDIT] Explicit parser metadata initialization (the uninitialized-meta warning)

Every compile emitted warning: `out parameter 'meta' may be uninitialized when 'IgParser' terminates`. The parser start state left **eight** classification fields unassigned on some terminal paths — `role`, `dir`, `fwd_port`, `port_ok`, `gen_in`, `dequeued`, `is_pktgen`, and `is_synth` (synthetic builds) — relying on the compiler’s implicit zero-init. The behaviour was safe (an unassigned `port_ok` defaults to 0, i.e. *fail-closed*), but that is exactly what the compiler declines to prove, so the warning stood on every build.

The repair assigns those eight fields **explicitly** in `start`, at the same zero encodings the code already relied on. Because each value equals the prior implicit zero-init, the compiled match-action pipeline is **bit-identical** — no table, stage or placement change — and the fix uses **no suppression pragma**. On the deploy compiler (bf-p4c 9.13.2): **0 errors and 0 uninitialized_out_param warnings across all four builds**, resource footprint unchanged. A probe first confirmed that a start-init-then-reassign is *not* a hard error on this compiler, contrary

to an earlier in-code assertion. Evidence: `evidence/final_silicon/`.

8 The state machine, and the bug that took two attempts

8.1 What the state has to express

One register byte, `reg_tag`, must express three things at once:

meaning	encoding	how it is distinguished
no transaction in progress	0x00	the value zero
live, no response yet	0xC0–0xCF	top bit set
live, response already queued	0x10–0x1F	top bit clear , never zero

The 0xC0–0xCF range is not arbitrary. DNP3’s application header byte for a first-and-final, solicited, unconfirmed request is 0xC in the top nibble and a 4-bit sequence number in the bottom — so *the protocol hands us* a per-transaction identity in exactly that range, and the parser accepts a READ only if its byte is in it. That is also the proof that **no live identity can ever be zero**: there is no arithmetic on it anywhere, the sequence advances inside the low nibble only, and the top nibble is pinned by the parser’s own acceptance test.

The third state is reached by **adding 0x50**: $0xCn + 0x50 = 0x1n$. That is one hardware instruction; it is one-shot for free, because adding 0x50 clears the top bit so applying it twice does nothing the second time; and it **keeps the identity**, since the low nibble still says which transaction. It is a state, not a flag. Figure 6 draws the result.

8.2 The bug: a missing exit

Gate 4 case C tested the obvious failure — the relay acknowledges but never answers. The result was clean on every count except one: **the transaction never ended**. `reg_tag` was left holding a live identity forever, and the *next* poll was therefore refused as a concurrent transaction: no reservoir, no hold, an unprotected transaction.

The cause is a two-line proof. There were exactly two ways to end a transaction: the released response ends it, and the fail-open budget ends it. With no response the first cannot fire; and the tokens die on the *deadline*, so they never reach their budget — the deadline **pre-empts** the second. Neither exit fires. The measured cost was exactly **one** subsequent unprotected transaction, after which the system self-heals, because that transaction’s own response clears the stale identity on its way out.

8.3 The repair that could not be built

The natural fix is to record “a response is pending for transaction *g*” in a second register, and on the ACK’s release check whether it matches. It **does not compile**: “*Table placement cannot make any more progress... dependency analysis has found that no more tables are placeable.*”

The reason is structural, not incidental. The response path must read the identity register *before* writing the pending register; the ACK-release path needs the opposite order. Register placement in this hardware is **static** — each register lives in exactly one pipeline stage — so two registers with opposite orderings on two paths form a dependency cycle, even though the two paths are mutually exclusive at run time. The condition that separates them is invisible to placement.

`p4/probe_retire_dependency.p4` reduces this to the smallest reproduction: two byte-sized registers, two paths, opposite orders, nothing else. `-DPROBE_CYCLE` reproduces the error verbatim; keeping the state *inside* the existing register (`-DPROBE_ONE_REG`) compiles in two stages. That is why the design is as it is.

8.4 The repair that was built

- when the first response of a transaction is admitted, add 0x50 (one-shot by construction);
- on the ACK’s release, **if the top bit is still set** — nothing pending — end the transaction immediately;
- if the top bit is clear, a response is queued, so leave the transaction live and let that response end it, exactly




Figure 6: The transaction state machine: three disjoint domains inside one byte. The two live states differ only in the top bit, so the switch separates them with a single signed comparison against zero (§7). The white lines inside the live boxes are what a circulating blocker token computes — the identity it carries minus the value it finds. Because the marking transition *adds* a constant rather than overwriting, that difference is 0x00 before marking and 0xB0 after, so one extra table entry covers all sixteen identities and the reservoir survives the state change. Source: figures/src/fig5_statemachine.py.

as before.

One extra table entry was required. When the marker changes, the circulating tokens are comparing themselves against a value that just moved. The difference a token computes is carried identity — stored value, which for a marked tag is $0xCn - (0xCn - 0xB0) = 0xB0$ for *every one* of the sixteen identities. So it is a single exact entry, not sixteen — and if it were wrong the failure would be immediate and loud: the tokens would read as foreign, the reservoir would collapse before the deadline, and the counters would show it. `BLOCK_TERM_STALE` was 0 in every subsequent test.

A free bonus fell out of the encoding. “Response pending” is a *distinct value*, not a flag, so every pre-existing test of the form “live and nothing pending yet” keeps its exact meaning — and a duplicate response therefore misses the hold branch by itself, with no new code. That is how §9 handles duplicates.

8.5 A defect that the repair itself introduced

Moving “idle” to 0x00 collided it with a *different* constant already meaning “leave this register alone” — also 0x00. Both transaction-ending paths write “idle” through the operation guarded by that constant, so **both silently became no-ops**. Nothing had run yet; it was caught by an audit, not by a failure.

General lesson: a fix that moves a sentinel value must enumerate every other sentinel in the same field. The distinct-value constant was moved to 0x01, which is safe because the field only ever holds that constant, idle, or an identity in 0xC0–0xCF.

8.6 The final state-transition table

event	before	action	after	packet's fate
READ, nothing in progress	0x00	write identity	0xCn	armed; mirror triggers 64 tokens
READ, one in progress	0xCn or 0x1n	no write	unchanged	refused as concurrent, forwarded unprotected
token, fresh	0xCn	read (increment 0)	unchanged	admitted to Q _{BLOCK}
token, difference 0x00 or 0xB0	either live state	read only	unchanged	still ours, loops again
token, foreign identity	any	read only	unchanged	stale, dropped
first response, in window	0xCn	add 0x50	0x1n	queued in Q _{HOLD} <i>behind</i> the ACK
duplicate response	0x1n	nothing	0x1n	suppressed
ACK release, response pending	0x1n	nothing	0x1n	forwarded; transaction stays live
ACK release, nothing pending	0xCn	write idle	0x00	forwarded; transaction ends here
queued response released	0x1n	write idle	0x00	forwarded; transaction ends here
response after the end	0x00	nothing	0x00	forwarded once, never held
tokens exhaust budget (fail-open, R2)	0xCn	note gen in 2nd reg	0xCn	token drops; reg_tag kept, not retired
next READ after a note (R2 recovery)	0xCn+note	arm on note or idle	new 0xCn	stuck slot reclaimed

The last two rows are the **repaired** fail-open (§7.7). The shipped design does **not** write idle on budget exhaustion — that destructive $0xCn \rightarrow 0x00$ write was defect 2 and is removed. A budget-zero token records its own generation in a second register (reg_failopen) and leaves reg_tag alone; the next READ consumes the note and re-arms only if reg_tag is idle or still carries the noted generation, so a fail-open recovers the slot instead of clobbering it.

analysis/test_tag_domain.py checks this model **exhaustively rather than by example**: all sixteen identities, all 240 ordered pairs of distinct identities, both markers, every transition. The E1 transition model alone is **2 256 assertions** (its pre-R2 count); with the R1 and R2 repair blocks the full suite is **2 674 assertions, 0 failures** today. The E1 core is mutation-checked four ways (revert the idle marker $\rightarrow$ 10 failures; re-collide the sentinels $\rightarrow$ 66; change the increment from $0x50$ to $0x40 \rightarrow$ 317; to $0x00 \rightarrow$ 195), because a test that cannot fail proves nothing.

9 Validation on synthetic traffic: gates 1 to 4

Before touching the relay, the mechanism was exercised with packets the switch generated itself. This is not a formality: it is the only way to test cases a real relay will not produce on demand, such as a response that never comes.

How the synthetic events are made, and a hardware law discovered doing it. The events (READ, ACK, RESPONSE) are emitted by the switch's own packet generator. The first attempt put all three in one generator run, spaced by a hardware gap. The reservoir then stood 1 000 012 ns after the READ instead of 1 215 ns — the ACK had already escaped.

The packet generator will not start a triggered burst while another generator run is in progress, and the wait equals the whole run's span — including its idle gaps. Measured at four points: a 3-packet run with a 200 μ s gap delayed the burst by 400 μ s; with a 500 μ s gap, by 1 000 μ s; two 1-packet runs with a 500 μ s inter-run gap, by 500 μ s; three with 200 μ s, by 400 μ s. In every case the tagged copy was emitted on time (688 ns) and the delay was *inside* the generator. Reproduced to the nanosecond: 1 000 012 ns.

Two further attempts failed for reasons worth knowing: a pattern-triggered generator application *cannot label its own packets* (all three decoded as the same one, so the roles collapsed), and two applications sharing one trigger are served events-first with no control-plane way to order them. The working schedule uses **two timer applications** — one emits the READ alone, the other the ACK and RESPONSE — **armed in a single write**, so the software skew (measured at 1.15 ms with two separate writes) cannot leak into the offset. The realised READ→ACK offset came out **10 ns** from target.

9.1 Gate 1 — does it load, and is the hardware configured

Both compilers agree exactly (9 stages, no drift); the program loads; strict priority is confirmed as $7 > 0$; $K = 64$ is confirmed; the fail-open horizon is confirmed as 30.802 ms.

On the *first* attempt this gate **aborted**, because the internal loopback port was at 10 Gbit/s instead of 25. That is not cosmetic: the token drain rate, and therefore both τ and H , scale with it. The check that caught it is now permanent and blocking.

9.2 Gate 2 — one transaction, seventeen requirements

PASS, 17/17. The headline is the decomposition already given in (5), with reservoir standing 678 ns, READ→ACK 500 010 ns, ACK→RESPONSE +28 ns (positive, so the order is right), and fail-open zero. Evidence: `evidence/gate2/gate2_20260729T231747Z/`, scored by `analysis/analyze_defense3.py`, whose self-test carries 17 negative controls proving each requirement can actually fail.

9.3 Gate 3 — five consecutive transactions, no reset between them

PASS, 5/5, 18 requirements each. The identity was *advanced* each time ($0 \times C0 \rightarrow 0 \times C4$), so a transaction that failed to end could not be mistaken for one that did — it would be refused as concurrent, exactly the signature described above.

quantity	min	max	spread
hold	2 001 427	2 001 586	159 ns
drain	1 691	1 695	4 ns
release tail	26	28	2 ns
reservoir standing	1 193	1 195	2 ns
READ→ACK	500 009	500 011	2 ns

The first attempt at this gate failed on our own criterion, not on the defense. The rule demanded that two registers be zero between transactions. They were not — they held the previous transaction’s deadline and identity — but the architecture never promised they would be: both are *self-clearing by design* (the next READ overwrites the deadline unconditionally; the response test is a *difference*, so a new identity reads non-zero with no reset). The replacement rule is **stricter**, not looser: it keeps “the transaction ended” and adds a check the old rule could not even see — that the inherited value must not collide with the new identity, which would silently invert the early/late classification.

9.4 Gate 4A — a response arriving just before the deadline

PASS 3/3, with the response arriving 4 872 ns before the deadline. Shrinking the margin from 1.5 ms to under $5 \mu\text{s}$ changed nothing, which is the point of the case.

9.5 Gate 4B — a response arriving after the acknowledgement has gone

PASS 3/3, response 500 128 ns late. It takes the normal forwarding path, forwarded exactly once, never held, and cannot disturb a later transaction. This is a *deliberate behaviour change*: before the state-machine repair a late response was briefly held; now the transaction has already ended, so it simply goes through. That is better — it costs one less internal loop.

9.6 Gate 4C — the relay never answers

PASS 3/3 after the repair, and this is the case that produced §8. The ACK’s release now ends the transaction, `reg_tag` returns to idle, and — the requirement that matters — **the very first following transaction is fully protected**, three times out of three. Before the repair, that transaction was completely unprotected.

9.7 A duplicate response, and an invariant that was being violated

If the relay retransmits its response while the first copy is still queued, what happens? The first answer was “the duplicate is forwarded” — and measurement showed that was wrong. Across three repetitions the duplicate departed 1 001 449, 1 001 341 and 1 001 421 ns *before* the held ACK.

The duplicate overtook the very packet the defense exists to delay, by 1.0014 ms, because the forwarding path goes straight out while the ACK is still queued. The gate had reported PASS: the rubric never tested ordering. It was found by adding a timestamp and looking.

The repair — the term for it is **current-transaction identity-matched response retransmission suppression** — drops such a copy while a response of that identity is already pending, and counts it separately. “Identity-matched” means: same TCP sequence position, same acknowledgement relationship, same learned session port, the same DNP3 solicited single-fragment framing, and the same transaction identity.

It is deliberately not called byte-exact. The response’s length and payload bytes are *not stored anywhere*, so they are not compared. A retransmission carrying the same sequence number but a *different* length is the one case this cannot distinguish. Storing the length would mean new permanent state, which the design does not have room for. Enqueuing a second copy instead of dropping it was rejected because a queued response ends a transaction unconditionally, so a second copy could end a *later* one.

After the repair: **3/3** — first response held and marking, duplicate suppressed, nothing forwarded early, and the bypass timestamp never written at all. A second measurement bug of ours was found the same way: the first version of that timestamp also fired on the *suppressed* copy, a packet that had been dropped and therefore departed nowhere.

9.8 A stale response arriving during a live transaction

The hardest isolation case. Transaction N finishes, $N+1$ arms with its reservoir standing and its deadline set, and then a response belonging to N arrives.

Making this test honest required work, because **the identity a response carries is its TCP sequence position**, and the ACK and the response are tested against the *same* register — so one packet template cannot produce “stale response, valid acknowledgement”. A third generator application with its *own* packet buffer and a sequence number offset by 0×1000 was added, firing $800 \mu\text{s}$ after the READ.

[AUDIT] THIS TEST’S PASS IS WITHDRAWN. It was reported as PASS 3/3 on the grounds that $N+1$ ’s identity, deadline, pending marker and 64 tokens were all unchanged and the stale copy took the bypass path. Two independent problems make that unsupportable.

First, the inference was backwards. The pending marker was said to be untouched “proven by the acknowledgement still finding the marker set”. That proves only that *some* response set the marker; it cannot say which — and §7.5’s defect 1 means the stale copy is itself able to set it before being rejected.

Second, the run’s own timestamps do not fit the intended schedule. The events were configured as READ +0.000 ms, ACK +0.500, **stale response** +0.800, legitimate response +1.000. The internal timestamps, identical to within ≈ 10 ns across all three repetitions, read:

<code>reg_ts_read</code>	+0.000 ms	
<code>reg_ts_ack_arm</code>	+0.500 ms	matches the configured ACK exactly
<code>reg_ts_resp_bypass</code>	+1.000 ms	the legitimate response’s slot, not the stale one’s

The bypass timestamp is written only by the arm that actually forwards, and only one bypass occurred, so it is a single unambiguous write — and it lands $200 \mu\text{s}$ from where the stale copy was scheduled. Compounding this,

the harness reads back the counters of the blocker generator and of applications 2 and 3 **but never application 4**, the stale injector, so there is no record of when it fired or which response the switch held.

Either the injector fired $200\ \mu\text{s}$ late, or the legitimate response was the one bypassed and the stale copy was held — an inverted test. **The evidence cannot tell these apart, so the case is recorded as UNRESOLVED, not passed.**

What the counters *do* say, in all three repetitions, is `RESP_HOLD_EARLY = 1`, `RESP_BYPASS = 1`, `RESP_DUP_SUPP = 0`. That pattern rules out the failure chain the audit predicted — the stale copy marking and the legitimate response then being suppressed — but it does not establish the intended behaviour either.

RESOLVED on the repaired build, 30 July 2026. The rerun was done and the case passes — on external evidence rather than inference. Nothing inside the chip could ever separate the two RESPONSES, so **the stale injector was given its own ethertype** (`0x88C8` against `0x88C7`), which costs no state and compiles free. The property then reduces to a sign: a bypassed copy is forwarded at once, a held copy waits for the deadline. From the master-side capture, **the stale copy left 1.514 ms BEFORE the held ACK in 6 of 6 repetitions** (min 1.431, max 1.530). It took the bypass path; $N+1$'s own RESPONSE stayed behind the ACK. The internal timestamp reconciles exactly: the stale copy arrives at `READ+1.000 ms` and bypasses at once, the ACK is released at `READ+2.501 ms`, and the 1.501 ms difference matches the wire.

And the thing that was actually wrong was the check. The withdrawn version asserted that the bypass timestamp equals the injector's *configured* offset. It does not, because **app 4's one-shot timer does not fire where it is configured** — $600\ \mu\text{s}$ and $800\ \mu\text{s}$ both realise at `READ+ $\approx$ 1000  $\mu\text{s}$` . That harness-fidelity defect produced the original $200\ \mu\text{s}$ discrepancy behind the withdrawal. It is now an explicit INFO line so it stays visible, and the case still exercises the intended condition because the realised arrival is well inside the hold window. Scored by `analysis/analyze_capture_f.py`, which carries four negative controls — a stale frame arriving *with* the ACK fails, an empty capture is INDETERMINATE rather than PASS. Detail: `evidence/repaired/RESULTS.md`.

9.9 Resource cost of the whole thing

There are **two generations** of the build, and they must not be conflated: the *original campaign* build (before the audit repairs) and the *final R1+R2+R3* build. The original D-sweep (§10–§11) ran on the first; the repaired campaigns (§10.5) and every §7 repair result on the second.

build generation	core	+ telem.	synth./inj.	egress	crit. path	errors
original campaign (unrepaired)	9 / 12	10 / 12	9 / 12	0	8	0
final R1 + R2 + R3	10 / 12	11 / 12	11 / 12	0	10	0

The state-machine repair of §8.4 cost **zero** stages (an intermediate version cost one — a write-after-write on a single metadata field; collapsing it to one write recovered both the stage and the critical path). The jump from the original 9/12-at-path-8 to the final 10/12-at-path-10 is **R1's** cost: authorising the marker before writing it adds one table and one dependency level, and stage count now equals critical path, so the final program is dependency-bound at 10. R2 and R3 add **zero** on top of R1 (§7.6).

[AUDIT] Which build produced the physical results. Every physical number in the original D-sweep (§10–§11) was collected on the **original 10-stage instrumented build**, because the hold decomposition needs `reg_ts_last_block` and `reg_ts_last_term`, which exist only under `D3_LIVE_FULL_TELEMETRY`; the 9-stage core of that generation has a compile and resource result only. The **repaired campaigns of §10.5** ran on the **final 11-stage R1+R2+R3 build**. The added telemetry registers are write-only and the critical path was unchanged by them, which supports functional similarity — but on a timing system that is an argument, not a proof, and a physical parity run on the final core build is open work.

10 Validation on the real relay

10.1 The physical setup, read from the hardware

The topology of Figure 1 was read out of the switch’s live configuration rather than assumed from file names — which mattered:

role	port	front panel	link
internal loopback, where the tokens circulate	8	15/0	25 Gbit/s, up, MAC-near loopback
the master (a Linux host, 192.168.10.1)	9	15/1	25 Gbit/s, up
the SEL-751 relay (192.168.10.7)	64	33/0	1 Gbit/s, up
replay injector (deliberately absent)	11	—	absent

On a freshly loaded program, none of this exists. Before the control plane ran, the port table had *no entry at all* for ports 8, 9 or 64, the traffic manager believed every port was 10 Gbit/s, and **all three loopback queues were at low priority** — no strict-priority ladder whatsoever. A test started in that state would have had no reservoir priority and would have silently measured nothing. Reading the state instead of trusting the configuration file is what caught it.

10.2 Stage 2 — the connection only, no request

A TCP connection was opened to the relay and held for 25 seconds with **zero bytes sent**, to check two things: that the switch learns the session correctly from real traffic, and that ordinary non-transaction acknowledgements do not accidentally trigger the defense.

Session learning was exact. The switch learned the master’s port number (32997) and the expected relay sequence position, both matching the packet capture precisely. The expected-acknowledgement register stayed *zero*, which is correct because only a READ installs it, and none was sent. **Nothing armed:** no identity, no deadline, no tokens, no hold, no queue drops.

The keepalive guard was exercised for the first time. The relay emits a bare acknowledgement roughly every ten seconds to keep the connection alive. The capture caught two, at +10.004 s and +20.024 s, plus one more when the connection closed. All three were **rejected** and forwarded normally; none armed a deadline, none entered the hold queue.

This matters more than it sounds. These keepalives look exactly like the packet the defense wants to hold: same source, same size, same flags, no payload. The conjunct that separates them is the TCP *sequence* position — a keepalive re-acknowledges data already acknowledged, so its sequence does not match what the switch expects. Earlier analysis over 8 captures had shown that the *acknowledgement*-based test accepts 61 of 61 keepalives while the *sequence*-based test rejects 61 of 61. The synthetic test build cannot produce a keepalive at all, so this guard had never been tested until this moment.

10.3 Stage 3, first attempt — a real failure worth recording

One Class-0 READ was sent. The frame was constructed and verified *before* it touched the relay: 18 bytes, matching the length the switch expected; addressed from master 1 to outstation 0; application byte 0xC0; function code 0x01 (READ); object group 60 variation 1, “all objects” — Class 0, read-only.

A real transaction completed: 134-byte response, healthy TCP. And the acknowledgement was **not held**: it reached the master 0.480 ms after the READ, when it should have been about 2 ms.

The registers gave the answer immediately: the packet generator was **not enabled**. `app_enable = false`, zero triggers, zero packets, zero tokens admitted. The synthetic test driver enables the generator itself as part of arming each transaction — the *live* control path had no equivalent step, and the mandatory cleanup at the end of configuration disabled it again.

This was a control-plane omission, not a mechanism defect: with an empty Q_{BLOCK} , every observed value was exactly what the design predicts. The run was **stopped and preserved** rather than patched over by increasing D or delaying the relay.

Everything else in that run worked, on real traffic, and each item was the first physical evidence of itself: the

real DNP3 READ armed a transaction through the live parse chain; the mirrored copy returned; **the real relay acknowledgement passed every acceptance test including the sequence conjunct**; the deadline armed exactly once at $t_{\text{ACK}} + D$; the state-machine repair fired and ended the transaction; and the response was forwarded exactly once.

10.4 Stage 3, second attempt — the hold works

With the generator armed and *left* armed — “armed” is a standing condition in the live path, because the trigger is the master’s own READ, which can arrive at any time:

quantity	value
first token admitted	2 769 950 513
all 64 tokens admitted	2 769 951 029 = 516 ns after the first
the relay’s real acknowledgement arrives	2 770 391 734
reservoir standing before the real ACK	441 221 ns = 0.441 ms
deadline armed	$t_{\text{ACK}} + 1\,999\,691$ ns, within quantization
first token notices the deadline	6 ns after it passed
last token gone	drain 1 693 ns
acknowledgement leaves	release tail 26 ns
hold	2 001 415 ns — corrected error – 168 ns
identity at the release	0×10 — the pending marker
identity afterwards	0×00 — ended by the queued response

Every Stage 3 requirement was met: exactly one trigger and 64 tokens, no admission drops, no duplicate hold, deadline armed once, no acknowledgement before the deadline, the response queued *behind* the acknowledgement, all 64 tokens terminating on the deadline, no stale terminations, no fail-open, transaction ended, no queue drops, TCP healthy.

And on the wire, which is what an observer sees:

event	reservoir armed	reservoir disarmed
relay acknowledgement reaches the master	+2.548 ms	+0.480 ms
relay response reaches the master	+2.590 ms	+5.015 ms
gap between them	+42 μ s	+4.535 ms

The relay emits its acknowledgement at about +0.48 ms either way. Armed, it arrives at the master at +2.548 ms — about 2.07 ms of added delay, matching the switch’s internal 2 001 415 ns. The response follows only 42 μ s later because it was queued behind the acknowledgement rather than travelling independently.

10.5 [AUDIT] The repaired build against the same relay

Everything above was measured on the build carrying both state-ordering defects. After R1 and R3 were repaired, the **live** repaired build was run against the same relay under the same campaign design — same six arms, same interleaving, same 200 ms gap, same D values, with polls per block doubled to 40. **960 attempted, 960 responded, 0 unanswered.**

The hold is unchanged on the wire. READ→ACK median, repaired against original: 1.519/1.514 at $D=1$, 2.517/2.515 at $D=2$, 4.514/4.508 at $D=4$, 8.587/8.519 at $D=8$ and 16.510/16.509 ms at $D=16$ — differences of 1–6 μ s against holds of 1–16 ms.

The CLRT result reproduces. At $D = 16$ ms: median 0.031 ms, sd 0.011 ms, max 0.049 ms, **160/160 collapsed**, against 0.032/0.012/0.047 and 80/80 originally. Still 22 distinct values, so still a distribution rather than a constant.

The mechanism stayed clean over 800 defended transactions: ordering invariant **960/960**, admitted tokens +**51 200** = 800×64 **exactly**, all deadline-terminated, zero stale terminations, zero fail-open, zero duplicate

suppressions, zero queue drops. **The thin fail-open margin is confirmed by a second session:** $1.59\times$ at $D = 16$ here against $1.49\times$ before, so §6.3's original $8.8\times$ is now wrong on two independent campaigns.

Two things this does not show. This session's relay was noisier — native CLRT sd 3.504 ms against 2.854, drift floor 0.582 against 0.530 — so the *between-session* separability differences must not be attributed to R1; that is what interleaving the arms guards against, and no claim rests on them. And the relay never sent a mis-sequenced response, so **R1's rejecting arm never fired:** this campaign establishes that the repair does no harm on the live path, not that it does good there.

A third campaign then added R2, same design, 960 more transactions. READ→ACK medians land within $1\text{--}6\ \mu\text{s}$ of the *unrepaired* original at every D — 1.513, 2.514, 4.514, 8.519 and 16.512 ms — and the CLRT result reproduces a third time ($D = 16$: median $32\ \mu\text{s}$, sd $13\ \mu\text{s}$, 160/160 collapsed). Ordering held 960/960, tokens were $+51\ 200 = 800 \times 64$ exactly, and stale terminations, budget expiries, duplicate suppressions and queue drops were all zero.

That run also **settles a loose end:** the R1+R3 session showed $D = 8$ at 8.587 ms, $68\ \mu\text{s}$ high, which I attributed to session noise rather than to R1. With a cleaner session — drift floor 0.511, the lowest of the three — it returns to **8.519 ms, matching the original exactly**. The attribution was right, and it is now evidence rather than an assertion. **R2's own path was not exercised on the live path either**, by construction: fail-open needs the budget to expire before the deadline, and in a healthy campaign the deadline always wins.

11 The D -sweep campaign, the data, and the analysis

One transaction proves a mechanism. It says nothing about whether the defense conceals anything. That needs a campaign.

11.1 How it was designed, and why

480 attempted transactions: 6 arms $\times$ 4 rounds $\times$ 20 polls, one TCP connection per block, 200 ms between polls, D changed at run time with no reload. Four design decisions, each forced by a known hazard:

Arms interleaved round by round, not run one after another.

Earlier sessions with this relay produced native-versus-native separability up to 0.985 — meaning two recordings of the *same undefended device* on different days can look almost as different as defended versus undefended. Session drift exceeds the effect, so every arm appears in every round and every comparison is made inside one session.

$D = 1$ ms is a pre-registered sub-threshold arm.

It is below the native CLRT, so theory says it should collapse nothing, and declaring that in advance makes it a check on the pipeline rather than a result. **[AUDIT] It was previously called a “null control”, which is wrong:** a null arm has no treatment effect, and $D = 1$ has exactly the effect the model predicts — it shifts the CLRT median by about 1 ms ($2.828 \rightarrow 1.799$). It is a low-dose arm; **the true null is the native arm**.

Attempted transactions counted, not successful ones,

with the disposition of every one. Counting only the ones that worked hides failures.

A separability number reported beside every concealment number.

Concealment alone is half a result.

Separability is the area under the ROC curve — the probability that a randomly chosen defended transaction and a randomly chosen undefended one can be told apart by that one number — folded onto $[0.5, 1.0]$, since an adversary may invert its own rule. Formally, for defended sample X (n values) and undefended sample Y (m values):

$$A = \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m \left[\mathbf{1}(x_i > y_j) + \frac{1}{2} \mathbf{1}(x_i = y_j) \right], \quad \text{sep} = \max(A, 1 - A). \quad (7)$$




Figure 7: (a) Every measured CLRT, one point per transaction, 80 per arm, on a logarithmic axis; black bars are medians and points are jittered horizontally only. The native cloud spans 1.7–13.2 ms; by $D = 16$ ms it has compressed onto an external floor of about 32μ s (median 0.0319 ms, sd 0.0120, max 0.0470 — a tight distribution, not a constant). Points in the grey band were measured as exactly zero, i.e. below the 1μ s resolution of the capture. **(b)** The same data plotted as the two observables against each other. The native cloud sits at the left, spread vertically — the vertical spread *is* the fingerprint. As D grows the cloud moves right and flattens: the spread leaves the CLRT axis and appears on the READ $\rightarrow$ ACK axis. Nothing is destroyed; it is moved. Source: `figures/src/fig7_scatter.py`.

sep = 0.5 is chance; 1.0 is perfect discrimination. It is computed on the **raw** measurement, with no binning, deliberately: binning has previously *raised* an entropy figure by 0.26 bits on a transform proven to be information-preserving, and a fully flattened feature can return perfect accuracy through density degeneracy rather than through information.

11.2 The result

480 attempted, **480 responded, 0 unanswered**. Native-versus-native separability — the drift floor — was 0.530, essentially chance, so this session was well behaved.

arm	D	CLRT med.	CLRT sd	CLRT max	collapsed	R $\rightarrow$ A med.	sep.
native	—	2.828	2.854	13.175	0/80	0.453	(floor 0.530)
d1 <i>sub-thr.</i>	1	1.799	3.331	15.465	0/80	1.514	0.649
d2	2	0.823	3.952	18.356	20/80	2.515	0.719
d4	4	0.032	1.129	7.888	63/80	4.508	0.966
d8	8	0.032	0.153	1.264	78/80	8.519	1.000
d16	16	0.032	0.012	0.047	80/80	16.509	1.000

All times in milliseconds; “collapsed” means the observed CLRT fell below 0.1 ms. **The hold tracks D exactly:** READ $\rightarrow$ ACK median is $D + 0.51$ ms at every single D . **The ordering invariant held in 480 of 480 transactions** — the acknowledgement was committed before the response, every time.

Figure 7 shows the raw measurements rather than summaries: every CLRT that was measured, native and defended.

Is the CLRT distribution compressed? Yes, by a factor of about 238. At $D = 16$ ms the standard deviation falls from 2.854 ms to 0.012 ms.

[AUDIT] It is not flattened to a constant, and an earlier version of this box said it was. The measured $D = 16$ sample: $n = 80$, **18 distinct CLRT values**, median 0.0319 ms, sd 0.0120 ms, minimum 0.0000 and maximum **0.0470** ms, with only **29 of 80** within $\pm 0.5 \mu$ s of the median and 10 at or below the 1μ s capture resolution. “All




Figure 8: The central result, 480 transactions against the real relay. (a) The observed CLRT distribution per arm. The whole distribution collapses onto the external floor — the dashed line — as D grows past the relay’s own response time. (b) The same sweep read two ways: the percentage of transactions whose CLRT is concealed (green), and how well an adversary separates protected from unprotected traffic using the CLRT (orange) or using READ→ACK (red). The dotted line is the native-versus-native drift floor, 53 %, which is what “no information” looks like in this session. **Concealment and detectability rise together, and detection is already near-perfect where concealment is only partial.** Source: figures/src/fig1_dsweep.py.

80 land on the same $32 \mu\text{s}$ constant” was an overclaim contradicted by this report’s own table. Compression is not equality.

[AUDIT] And “concealed” is the wrong word for the count in the table. “Collapsed” there means one thing only: the observed CLRT fell below an 0.1 ms threshold. That threshold is a choice, and clearing it does not make the device unidentifiable. At $D = 4$ ms, 63/80 clear it while the CLRT still rank-separates from native at 0.966 — the feature has not been concealed, it has been transformed into a different, highly recognizable distribution. Read “**collapsed below threshold**” for the count, and reserve *concealment* for the question §12 says this campaign cannot answer.

11.3 The mechanism held up over the 400 defended transactions

Per-arm counter totals, identical for every armed arm:

measurement	value	meaning
tokens admitted	5 120 = 64×80	exactly K per transaction, every transaction
transactions armed	80	one per poll
tokens ending on the deadline	5 120	every token, the intended path
stale terminations	0	no token ever lost track of its transaction
budget expiries / fail-open	0	the safety net was never needed
admission drops	0	no token arrived to find no transaction
concurrent-transaction refusals	0	every transaction ended before the next began
duplicate suppressions	0	the relay retransmitted zero times in 480
queue drops	0	nothing was lost

And both state-machine exits partitioned exactly, 400 times:




Figure 9: Separability from native for each of the three timings an observer can measure, at every D . Bars start at 0.50, which is chance. **READ→ACK — the feature the defense creates — beats the CLRT at every single D ,** and the total READ→RESPONSE is the *least* separable, because while D is below the native CLRT the defense conserves the total and merely moves time from one term into the other. A bar at the dotted drift floor would mean no information at all. Source: `figures/src/fig3_observer.py`.

arm	response in window	arrived after	ended by the ACK	ended by the response
d1	0	80	80	0
d2	18	62	62	18
d4	62	18	18	62
d8	78	2	2	78
d16	79	1	1	79

The two columns always sum to 80, and “ended by the ACK” always equals “arrived after”. As D grows past the relay’s own response time, transactions migrate from one exit to the other. This is the state-machine repair sweeping across its own boundary on real traffic, 400 times, with no exceptions — the strongest evidence in this report that the state machine is right.

11.4 The second analysis, which changes how the first should be read

The table above scores *one* number. A real eavesdropper sees every timing the exchange produces — and Defense 3 **creates** one: READ→ACK becomes $D + 0.51$ ms.

Drift floors, native versus native, are at chance for all three features: READ→ACK 0.514, CLRT 0.530, READ→RESPONSE 0.503.

arm	D	READ→ACK	CLRT	READ→RESPONSE (the total)
d1 <i>sub-thr.</i>	1	0.898	0.649	0.542
d2	2	0.931	0.719	0.578
d4	4	1.000	0.966	0.669
d8	8	1.000	1.000	0.925
d16	16	1.000	1.000	1.000

And a classifier that is *not* fitted on the data it is scored on — a single threshold chosen from rounds 1–2 and tested on rounds 3–4, with balanced accuracy so class sizes cannot flatter it:

arm	D	threshold	balanced accuracy, held out
d1 <i>sub-thr.</i>	1	0.985 ms	0.863
d2	2	1.495 ms	0.950
d4	4	2.485 ms	0.963
d8	8	4.490 ms	1.000
d16	16	8.485 ms	1.000

[AUDIT] The 80 transactions in an arm are not 80 independent observations

They come from **four TCP connections of 20 polls each**. Polls inside one connection share the relay’s scheduler state, the connection’s state, host load and clock drift, so the effective replication is **4 blocks, not 80 transactions**, and every interval above is narrower than it should be. The single rounds-1–2 / rounds-3–4 split is better than scoring on the training data, but one split quantifies no uncertainty at all.

`analysis/analyze_blocked.py` redoes it properly: the bootstrap resamples **whole connections**, and the held-out score is **leave-one-round-out**. 4 000 block resamples, seed 20260730:

arm	D	READ→ACK sep. (95 % CI)	CLRT sep. (95 % CI)	LORO balanced accuracy
d1	1	0.898 (0.853–0.933)	0.648 (0.592–0.697)	0.906 (0.875–0.925)
d2	2	0.931 (0.895–0.959)	0.719 (0.668–0.763)	0.956 (0.925–0.975)
d4	4	1.000 (1.000–1.000)	0.966 (0.942–0.989)	1.000
d8	8	1.000 (1.000–1.000)	1.000 (1.000–1.000)	1.000
d16	16	1.000 (1.000–1.000)	1.000 (1.000–1.000)	1.000

Every point estimate survives, and **the detectability conclusion strengthens rather than weakens**: leave-one-round-out gives 0.906 at $D = 1$ ms where the single split gave 0.863, and no confidence interval at any D reaches down to the drift floor. The block-resampled drift floors are unchanged at 0.514 / 0.530 / 0.503. This still does not make the campaign a four-fold replicated experiment in the strong sense — four connections in one session against one device is what it is — but the widths are honest where the earlier point estimates were not.

Three things follow.

The CLRT number was the flattering one. READ→ACK beats it at every D , and the gap is widest exactly where the defense looked best: at $D = 2$ ms the CLRT is only partly concealed (20/80, separability 0.719) while one threshold on READ→ACK already gets 0.950 on data it never saw.

Detectability exceeds the collapse it buys, at every D . Even at $D = 1$ ms — the sub-threshold arm that collapses *nothing* — a leave-one-round-out classifier reaches 0.906. There is no setting in this sweep at which Defense 3 is harder to detect than the CLRT information it removes.

[AUDIT] One caveat on presentation. Figure 8(b) and (c) were previously a single panel plotting “% collapsed below 0.1 ms” and separability $\times 100$ on one percentage axis. Those are not the same kind of quantity — one is a thresholded sample proportion, the other a ranking statistic — so they must not be compared arithmetically, and the panel is now split to stop inviting it. The conclusion rests on the separability and the held-out classifier alone.

The reason is now measured, not argued. READ→RESPONSE, the *total*, is the **least** separable feature: 0.542 at $D = 1$ against a 0.503 floor — essentially unchanged. While D is below the native CLRT, Defense 3 approximately **conserves the total** and merely **redistributes** time out of the CLRT and into READ→ACK:

$$a + c \xrightarrow{\text{Defense 3, } D < c} \underbrace{(a + D)}_{\text{observable 1}} + \underbrace{(c - D)}_{\text{observable 2}} = a + c. \quad (8)$$

That is what “the leak is relocated, not destroyed” means, quantitatively: the observable that survives is the sum, and the leak is the redistribution. Only when D dominates the total ($D = 16$) does the sum separate too — and by then everything separates at 1.000.

11.5 What survives, precisely

READ→ACK after the hold is D plus the relay’s own acknowledgement latency. Since D is our own parameter, and therefore known to anyone who measures a few transactions, subtract it:

arm	R→A med.	R→A sd	med. $-D$	sep. of (R→A $-D$) vs native
native	0.453	0.825	—	(floor 0.514)
d1	1.514	0.798	0.514	0.690
d2	2.515	0.771	0.515	0.702
d4	4.508	0.680	0.508	0.679
d8	8.519	0.347	0.519	0.683
d16	16.509	0.588	0.509	0.660

The spread is essentially untouched — 0.35 to 0.80 ms after the defense against 0.825 ms before. The hold *translates* the acknowledgement-latency distribution by D ; it does not compress it. Subtracting D recovers something only 0.66–0.70 separable from native against a 0.514 floor: not a perfect reconstruction, thanks to a systematic +0.06 ms the switch adds, but far closer to native than the raw READ→ACK at 1.000.

So the relay’s own acknowledgement latency — a *different* device property from CLRT — passes through this defense with its shape intact.

12 What may and may not be claimed

12.1 Established

- The mechanism ran on real hardware.** The campaign contained **480 completed transactions, of which 400 were defended** (5 armed arms $\times$ 80) and 80 were the native arm with no reservoir and no hold. Across the 400 defended: exactly 64 admitted tokens each, 25 600 tokens in total, all terminating on the deadline, zero fail-open, zero queue drops. The acknowledgement-before-response ordering held in **480 of 480**. Hold accuracy –168 ns. *[AUDIT] Previously “480 of 480 transactions, exactly 64 tokens each” — impossible under the stated campaign, since the native arm has no tokens by construction.*
- The hold is governed by D and nothing else.** READ→ACK = $D + 0.51$ ms at five values of D spanning 1 to 16 ms.
- CLRT compression on this device.** Standard deviation $2.854 \rightarrow 0.012$ ms at $D = 16$ ms, a factor of about **238**; median $\approx 32 \mu\text{s}$, maximum $47 \mu\text{s}$, 18 distinct values. *[AUDIT] Previously “80 of 80 transactions flattened onto a $32 \mu\text{s}$ constant” — false.*
- The state model is exhaustively checked, and all three of the audit’s defects are now repaired.** The Python reference model passes **2 674** assertions and is mutation-checked, and the two physical exits partitioned exactly across 400 transactions. Of the three defects (§7.5): **R1** is validated on silicon and across 1 920 live transactions; **R2** on silicon, the fail-open path now crediting all 64 tokens to the budget instead of 1 (§7.7); **R3** on silicon, the forged frame dropped before it reaches the loopback (§7.8). Full compiled-state correctness is still not *proven* — the model is not the silicon — but no known defect remains unrepaired. *[AUDIT] Previously “the state machine is correct across its whole domain”.*
- Graceful degradation.** When D is smaller than the CLRT the output is $c - D$, not the untouched CLRT — a partial rather than a cliff-edge failure.
- The observed non-transaction traffic was not disturbed.** Three real relay keepalive acknowledgements were rejected and forwarded, plus 61 further captured examples used in offline predicate analysis. *[AUDIT] Previously a general claim from three physical observations.*
- Packets are not modified.** No byte of any forwarded packet is changed; only the time at which it leaves. This is unaffected by anything above.
- [AUDIT] The three repairs behave as designed on silicon.** R1’s authorisation table, R2’s fail-open note and R3’s injection drop were each exercised on the switch — R1 across 1 920 live transactions (1 600 de-

fended) and Gate 4 case F; R2 at two budgets and the K -sweep; R3 with the in-switch injector. Their behaviour against a *live* network adversary has limits, stated below.

12.2 Not established, and why

1. **Device anonymity.** Compressing CLRT is not the same as making the device unidentifiable, and this report does *not* demonstrate the latter. **[AUDIT] These are different classification problems and this campaign only measures the first:** *task A* is native SEL-751 versus defended SEL-751 — which is what §11 measures, and which is a *defense-detectability* result; *task B* is defended SEL-751 versus a defended relay of another model — which is what the threat model concerns, and which no data here touches. A high score on task A is a genuine secondary leakage finding, not a refutation of device concealment. The relay’s own acknowledgement latency survives with its spread intact. Worse, the question cannot be answered with this data at all: **there is no confusion set.** The other devices in the corpus are combined-ACK, separable on packet count alone. **One relay cannot answer a discrimination question**, however many transactions are run against it. Settling it needs a second *separate-ACK* device measured under the same D .
2. **“Better than Defense 2”.** No iso-latency Defense 2 arm was run — it needs a different switch program. And it must be matched on *added latency*, not on the parameter: comparing $D \leq 3$ ms against $G = 25$ ms is uninterpretable in both directions. The numerical target is now known, since Defense 3’s added latency is $\approx D$.
3. **Undetectability.** The opposite is established. Detection is perfect at $D \geq 8$ ms and 0.906 (leave-one-round-out) even at the sub-threshold arm.
4. **A steady-state CLRT for the SEL-751.** One relay, one session, one 200 ms poll rate. The relay has at least two timing regimes — a connection-cold first poll and a steady state — so the 2.828 ms median here is this session’s distribution, not the device’s.
5. $K = 64$ **is minimal.** It is sufficient, and it is what was measured. No smaller value was tested.
6. **Concurrency.** One active protected transaction at a time. This is the *measured capacity* of the mechanism — a hold consumes roughly 24 of the 25 Gbit/s internal loopback — not a prototype simplification that a later version removes.
7. **Segmentation.** Every response in the corpus and in every test was a single segment. Multi-segment responses are detected and forwarded unprotected, not handled.
8. **[AUDIT] Full compiled-state correctness is checked, not proven.** All three defects are repaired and each behaves as designed on silicon (§7.7, §7.8, §10.5), and the reference model passes 2 674 mutation-checked assertions — but the model is not the silicon, and no exhaustive proof over the compiled program exists.
9. **[AUDIT] The repairs against a real wire adversary.** R3 is demonstrated with an **in-switch** forged-frame injector (§7.8), not a frame from an external host on a real port — the lab has no such vector. R1’s rejecting arm is demonstrated **synthetically** (Gate 4 case F); on the live relay it never fired, and the topology has no host on the relay-facing port to forge one. The repairs are shown correct against the switch’s own generated traffic, not against a network attacker.
10. **[AUDIT] The cross-transaction clobber of defect 2 was never produced on hardware.** The *within*-transaction defect is fully reproduced: a single native budget-zero token clears `reg_tag` at $K = 1$, giving the $1/K - 1$ cascade (§7.8’s K -sweep), and R2 fixes it. The *cross*-transaction case needs the generation-wrap coincidence the harness cannot arrange, so it stays model-checked. The write it relies on is confirmed real and single-token; the cross-transaction *reach* is not.
11. **[AUDIT] The sub-microsecond retirement boundary.** Gate 4B placed the late response $500 \mu\text{s}$ after the acknowledgement’s release. The dangerous interval — after the acknowledgement has retired the transaction but before it has left the master-facing queue — was never tested. It needs a sweep at 0 / 32 / 64 / 128 / 256 / 512 ns / 1 μs measuring master-facing *egress* order, not *ingress* timestamps.
12. **[AUDIT] That the measurement point is the attacker’s wire view.** Captures were taken with `harness/block.py` **on the master host**, and a host PCAP timestamp is not a port-9 wire egress timestamp: send timestamps can precede transmission, receive timestamps follow reception, and the capture resolves to

about 1 μ s. The $\approx 32 \mu$ s floor may therefore be partly a capture artifact. Read every number as *measured at the master host interface, used as a proxy for the port-9 observer* — not as *exactly what the attacker gets*.

12.3 [AUDIT] What the implementation requires, which the earlier text did not disclose

None of these is a defect; all were undisclosed, and several decide whether a real substation deployment would be protected at all.

requirement	consequence if unmet
Plaintext DNP3 over TCP — the parser reads the function code, application control byte, transport FIR/FIN and TCP fields	end-to-end TLS or IPsec makes the exchange invisible to this parser. §1 argues correctly that encryption does not remove timing leakage, but this implementation cannot act on encrypted traffic ; it must sit at a plaintext point
Ethernet II, no VLAN tag — the parser transitions on EtherType 0x0800 directly	VLAN-tagged substation traffic bypasses the defense entirely
IPv4 with <code>ihl == 5</code>, <code>mf == 0</code>, <code>frag_offset == 0</code>	IP options or fragments bypass
TCP options ≤ 12 bytes on DNP3-bearing packets (<code>data_offset</code> 5–8; pure acknowledgements accept 5–15)	a response with more than 12 bytes of options bypasses unprotected
One configured TCP session, one active transaction	every state register has size 1, so the limit is one protected <i>connection</i> ; a second matching connection would overwrite the learned port and sequence trackers
Fixed configured READ payload length; any DNP3 READ matches	the parser checks <code>(app_control & 0xF0) == 0xC0</code> and <code>func == READ</code> only — no object group, variation or qualifier — so this is <i>evaluated using</i> Class-0 READs, not restricted to them
TCP sequence 0 sentinel — RESOLVED (§7.9)	earlier versions used seq 0 as a no-write sentinel, which declined to update after a sequence-space wrap; the final writer/reader split stores seq 0 unconditionally (assembly-verified, Gate 2 PASS). No longer a limitation
Duplicate suppression discards a retransmission	ordering is preserved, but if the queued original is later lost that recovery opportunity is gone; TCP recovers later
“Zero dropped packets” needs qualifying	the mechanism deliberately drops blocker tokens at the deadline, trigger clones, stale tokens, matching duplicates and off-topology frames. The defensible claims are zero queue drops and zero unintended host-packet drops

12.4 Open work, and what each item blocks

#	open item	what it blocks	why it is not done
1	iso-latency Defense 2 arm	any Defense 2 vs Defense 3 statement, either direction	needs a different switch program loaded; G must match <i>added latency</i> , so $G \approx D + \text{native CLRT}$
2	a D calibrated on one campaign and tested on another	selecting an operating point	the governing constraints forbid fitting and testing D on the same campaign; the sweep here does not fit, so nothing is violated — but nothing is selected either
3	a second separate-ACK device	the device-anonymity question in any form	not available; a corpus limitation, not a schedule one
4	safety tests as a named stage	nothing already covered	fail-open, keepalive, concurrent, stale and duplicate cases are all exercised above; superseded in substance, never in form

#	open item	what it blocks	why it is not done
5	the $D = 40$ ms clamp (was INFEASIBLE) — DONE	—	the fixed 40 ms clamp is removed ; <code>control/parameter_policy.py</code> computes $D_{\max} = H - a_{\text{bound}} - t_{\text{detect}} - t_{\text{drain}} - t_{\text{tail}} - M \approx 24.8$ ms and refuses $D = 40$ ms (verified on silicon, §6.3). The supported range is now correct by construction
6	multi-segment responses	nothing claimed	every response in the corpus was a single segment, so the path has never been taken
7	rollback to Defense 2	nothing; it is deliberate	the switch is intentionally left running Defense 3 with the reservoir armed
8a	repair defect 1 — DONE (R1)	—	repaired, validated on silicon in the synthetic build (§7.6) and against the physical relay over 960 transactions (§10.5). Remaining: an adversarial live case that actually presents a mis-sequenced response
8b	repair defect 2 — DONE, validated on silicon	—	second-register design, free on top of R1+R3, assembly-asserted, model-checked and confirmed on hardware at two budgets (§7.7). Remaining: a FOREIGN token reaching budget zero while a later transaction is live — the case the defect was dangerous in, which the harness cannot currently arrange. (The live build is done — the final R1+R2+R3 campaign ran, §10.5.)
9	remove the host-injected <code>0x88C1</code> path — DONE, demonstrated (R3)	—	closed in source, and an in-switch injector shows the forged frame dropped at the fresh stage under R3 and entering without it (§7.8)
10	[AUDIT] eliminate the uninitialized-metadata warning	nothing observed — if the metadata really is zeroed the default is <code>port_ok = 0</code> , i.e. <i>fail-closed</i> — but that is exactly what the compiler declines to prove	present in every build log; assign every load-bearing field on every terminal parser path
11	rerun §9.8 with the stale injector identifiable — DONE	—	resolved on the repaired build with master-side capture, 6/6 (§9.8). Remaining: wrong-port and wrong-acknowledgement variants, and the app-4 timer defect — its one-shot fires at $\approx 1000 \mu\text{s}$ regardless of the configured offset
12	[AUDIT] sweep the retirement boundary at 0–1 μs	the narrowest ordering guarantee	the master-facing ACK-vs-RESPONSE order is measurable with the hardware RX timestamps of #14 (Vision’s NIC), so this is a ready experiment , not a hard block
13	[AUDIT] a physical parity run of the core build against the instrumented build	that the core behaves as the telemetry build the timing is inferred from	artifact-level parity established (core vs telemetry 9.13.2 assemblies share bit-identical SALU logic, differing only by the two write-only timestamp registers); a full physical core-build campaign remains the open part
14	[AUDIT] hardware-timestamped capture	that the $\approx 32 \mu\text{s}$ floor is a wire property, not a capture artifact	now known ACHIEVABLE : Vision’s capture NIC supports hardware RX timestamps (<code>ethtool -T: hardware-receive/raw-clock</code>), resolving this at ns rather than the $\approx 1 \mu\text{s}$ software floor — a ready experiment

#	open item	what it blocks	why it is not done
15	a control-plane guard on the poll rate — DONE	—	implemented in <code>control/parameter_policy.py</code> : it enforces $16 T_{\text{poll,min}} > H + t_{\text{drain}} + M$ and refuses a too-fast poll (a 2 ms interval is rejected by its own self-test). The remaining part is optional: physically reproducing the cross-transaction generation-wrap coincidence

12.5 [AUDIT] Final status — what is closed and what remains open

Stated as one line: **all audit remediation items are complete** — the three behavioural defects (R1/R2/R3), the parser-metadata warning (§7.10) and the TCP-sequence-zero repair (§7.9) are resolved and validated on silicon. **Several broader evaluation questions remain open** — external-wire adversarial injection, cross-transaction generation-wrap reproduction, a hardware-timestamped observer capture, the ACK-retirement boundary sweep, K-minimization, and a full physical core-vs-telemetry parity campaign — but these are not audit-remediation items and not release gates.

area	status
R1 — response authorisation	closed — synthetic rejection (Gate 4 case F) + live non-regression over 960 txns
R2 — fail-open accounting / recovery	closed for the demonstrated <i>within-transaction</i> defect (K TMO / 0 STALE, <code>reg_tag</code> preserved; K-sweep reproduces the write at K=1)
R2 — cross-transaction generation-wrap reach	model-checked, not physically reproduced
R3 — fresh external-token admission branch	closed on silicon via the in-switch injector (dropped fresh, counted BLOCK_REJECT)
BLOCK_ENQ / BLOCK_REJECT accounting	closed and re-verified on silicon
broad relay campaigns	sufficient — do not repeat
external <i>wire</i> adversary (R1/R3 from a real host port)	topology-blocked — dp64 faces the SEL-751 directly, no host to inject from (§12.2)
acknowledgement-retirement egress sweep (0–1 μ s)	ready experiment — measurable via Vision’s hardware RX timestamps (open-work #12)
hardware-timestamped observer capture	ready experiment — Vision’s NIC supports hardware RX timestamps (open-work #14)
core-vs-telemetry parity	done at artifact level ; full physical campaign open (open-work #13)
parser uninitialized-meta compiler warning	RESOLVED (§7.10: parser start-init; 0 warnings on all four 9.13.2 builds)
TCP-sequence-zero sentinel (§7.9)	RESOLVED — writer/reader split, seq 0 now stored (assembly-verified, Gate 2 PASS)
documentation / artifact consistency	reconciled 2026-07-31 across report, README, P4 header, resource ledger, MANIFEST, and all counts/totals

12.6 Stated head-on rather than buried

At the values of D that actually collapse the CLRT, **the output is physically implausible**. A device that took 16.5 ms to acknowledge a 20-byte read and then answered 32 μ s later, every single time, resembles no real relay: generating an acknowledgement is cheap and computing an answer is expensive, so the defense **inverts the natural ordering of costs**. Defense 2’s output stays inside the space of plausible device behaviour; Defense 3’s, at the D that conceals, leaves it.

Safety. Every physical transaction in this report was a *read*. No SELECT, no OPERATE, no DIRECT OPERATE, no configuration write, no setting change and no cold restart was ever sent to the relay — not once in the **2 400 transactions across the three physical campaigns** (480 original + 960 + 960) plus the smoke tests. The defense never modifies a byte of any packet; it only changes *when* packets leave.

13 How to reproduce everything

13.1 Software only, no hardware

command (from defense3/)	what it checks
python3 analysis/test_tag_domain.py	2 674 assertions; exit 0
python3 analysis/analyze_defense3.py --self-test	17 negative controls
python3 analysis/analyze_gate34.py --self-test	20 controls
python3 analysis/analyze_check2.py --self-test	6 controls
python3 analysis/analyze_dsweep.py evidence/physical/dsweep_blocks.jsonl /tmp/a.json	regenerates the sweep tables
python3 analysis/analyze_observer.py evidence/physical/dsweep_blocks.jsonl /tmp/b.json	regenerates the observer tables

The last two regenerate every table in §11 from the raw per-transaction data.

13.2 Compiling

```
bf-p4c --target tofino --arch tna -g \
[-DD3_SYNTH_EVENTS | -DD3_LIVE_FULL_TELEMETRY | -DD3_INJECT] \
-o <outdir> \
p4/case_a_defense3.p4
python3 analysis/assert_salu_asm.py <outdir>/pipe/*.bfa # this MUST pass
```

Compile the **canonical** p4/case_a_defense3.p4 (R1/R2/R3 unconditional — a no-flag build is the safe repaired program). Final footprints: no flag = core, **10/12 stages, critical path 10**; D3_LIVE_FULL_TELEMETRY = core plus the two internal timestamps, **11/12, path 10**; D3_SYNTH_EVENTS = the synthetic gate build, **11/12, path 10**; add D3_INJECT for the adversarial injector (synthetic only). Never compile the synthetic flag for live use — it relaxes an acceptance conjunct so that generated packets can reach the real hold path. (The pre-audit unrepaired build under archive/pre_audit/ is 9/12 at path 8: the historical control, not production.)

13.3 The figures

Every figure in this document is regenerated by one script under figures/src/, run with \$RESEARCH_PYTHON:

script	figure
fig1_dsweep.py	Fig. 8 — the <i>D</i> -sweep result, three panels
fig2_mechanism.py	Fig. 3 — construction and hold decomposition
fig3_observer.py	Fig. 9 — per-feature separability
fig4_timelines.py	Fig. 2 — the four defenses on one axis
fig5_statemachine.py	Fig. 6 — the transaction state machine
fig6_trigger.py	Fig. 4 — the trigger chain and its margin
fig7_scatter.py	Fig. 7 — every raw CLRT, native and defended
fig8_topology.py	Fig. 1 — the physical setup
fig9_ksweep.py	Fig. 5 — the fail-open <i>K</i> -sweep reconciliation

Each script reads evidence/physical/dsweep_blocks.jsonl or the measured constants quoted here, re-computes every number it plots, and prints them, so a figure can be checked against the tables. Output is vector PDF plus 300 dpi PNG at IEEE column widths (3.5 in single, 7.16 in double) with 9 pt Times New Roman, so nothing is rescaled on the page. Palette alessandretti-nature, one colour per meaning across all nine

figures. `analysis/analyze_blocked.py` regenerates the block-clustered intervals in §11; it carries seven negative controls of its own.

The one deviation from the figure conventions: the schematics are drawn in matplotlib rather than Inkscape. That trades a little typographic polish for the diagrams being *regenerable from the same scripts as the data panels* — no manual step between the measurements and the drawing.

13.4 On hardware

Loading displaces whatever is running and needs explicit authorization. The synthetic gates are `./run/run_defense3.sh --gate2` (one transaction), `--gate3` (five consecutive), `--gate4` (the boundary cases) and `--check2` (trigger latency, 100 trials). The runner loads nothing itself, asserts the loopback speed before and after, refuses to start on dirty state, and always restores — deliberately delegating restoration to the one existing copy of that code rather than reimplementing it.

The physical campaign is `harness/campaign.sh out.jsonl 4 20 0.2` (rounds, polls per block, gap in seconds), with `harness/setarm.py` on the switch (sets D , arms the reservoir, clears per-transaction state) and `harness/block.py` on the master (captures, polls, parses the capture into per-transaction rows).

14 Every mistake made, and what it cost

Recorded because the pattern is more useful than the individual bugs: **in this project, tests and criteria were wrong about as often as the code was.**

#	mistake	how it was found	cost
1	Graded the design against a broader objective than the threat model sets, and concluded “do not build”	corrected by the project lead	one wrong verdict, reversed
2	Quoted $D = 12$ ms as reaching a 99th percentile of 12.607 ms	caught in review	corrected to 13
3	Pooled a connection-cold poll into a “steady-state” sample	a review pass	D for a full clamp was 13 ms, not 22
4	Large constant in stateful hardware; the write silently never committed	reading the compiled assembly, after two wrong theories	the whole of §7
5	Then claimed the <i>cause</i> was proven from the assembly	a 13-constant probe showed identical output for all K	claim narrowed to an inference
6	Moving the idle marker collided it with a different sentinel; both transaction exits became no-ops	an audit, before it ever ran	would have broken every second transaction
7	Counted the defense’s own mirrored copy as an off-topology packet	the same audit	made a gate requirement unsatisfiable whenever the defense worked
8	Left a stale <code>0xFF</code> in the scoring code	the same audit	a gate could never pass again
9	Put all synthetic events in one generator run	measured 1 000 012 ns, reproduced to the nanosecond	discovered the generator run-span law
10	Assumed a pattern-triggered generator can label its packets	three roles collapsed into one	forced the two-timer design
11	Set the token increment in the wrong branch of the classifier	16 admitted then 48 dropped: $0xC0 + 16 = 0xD0$ leaves the valid range exactly at token 17	one silicon run
12	Clean-state criterion demanded zeros the architecture never promised	Gate 3 stopped at transaction 2	replaced with a <i>stricter</i> rule
13	Sign test written unsigned; always false	the assembly assertion, which now blocks it	§7
14	Duplicate-response rubric never tested ordering	added a timestamp and looked	the duplicate was overtaking the held ACK by 1.0014 ms

#	mistake	how it was found	cost
15	That new timestamp also fired on the <i>dropped</i> copy	the gate failed against a packet that departed nowhere	one analysis pass
16	Scored a boundary case with the normal-transaction rubric, which forbids the bypass the case exists to produce	the gate failed on its own purpose	one analysis pass
17	No live arming step, so the first physical run had an empty hold queue	the registers said <code>app_enable = false</code>	one physical run, stopped and preserved
18	Per-block counter zeroing failed silently; cumulative snapshots were summed	the totals were absurd (307 arms for 80 polls)	one arm's counters unusable; recovered by differencing
19	Reported concealment on one feature	asked what an eavesdropper actually gets	§11 — the headline result changed

Found by the external audit of 30 July 2026:

20	Wrote state before validating it, twice — the response marker and the fail-open retire both commit at pipeline level 2 while their authorising test resolves at level 3	an external audit read the source	§7.5; two confirmed unfixed defects, and it invalidates §9.8
21	Called a stale-response case PASS on an inference that could not distinguish the two responses , and the run's own timestamps put the single bypass 200 μ s from where the stale copy was scheduled	the audit, then re-reading my own evidence file	§9.8 withdrawn
22	Sized the fail-open horizon against D alone , quoting $8.8\times$ from a stale design point and omitting the relay's own ACK latency	the audit	the true worst case was $1.49\times$, and the advertised 40 ms clamp is infeasible
23	Said “480 of 480 transactions, exactly 64 tokens each” when only 400 were defended	the audit; my own exit-partition table already totalled 400	an internal contradiction that survived every read
24	Said 80 transactions “land on the same 32 μs constant” when the sample has 18 distinct values and a 47 μ s maximum	the audit, contradicted by my own table on the same page	compression restated as compression
25	Used “release tail” for two quantities three orders of magnitude apart	the audit	§6.2 now names them separately
26	Said the compiler “accepted all four traps without complaint” in a section whose third trap is a hard compiler error	the audit	§7 reclassified; the unsigned comparison is a type error, not a miscompile
27	Treated 80 transactions as 80 independent observations when they came from 4 connections	the audit	§11 now block-bootstraps by connection; the conclusion strengthened
28	Called $D = 1$ ms a null control when it produces the predicted 1 ms shift	the audit	relabelled a sub-threshold arm; the native arm is the null

Found while repairing the defects the audit identified:

29	The repair for defect 1 broke the mechanism: its authorisation table used a catch-all default setting <code>tag_val = 0</code> , which for every non-RESPONSE class turned the tag write into an unconditional <code>TAG_INACTIVE</code>	Gate 2, on the first transaction, on silicon	the trigger clone wiped the generation ≈ 700 ns after the READ armed it. It had already passed 2 354 offline assertions and a compile-fit check
30	Asserted the stale injector arrives where it is configured. It does not — 600 and 800 μ s both realise at ≈ 1000 μ s	the master-side capture, after the check failed 6/6 and the mechanism turned out to be right	the check was scoring harness fidelity, not switch behaviour
31	Said the capture could not label which frame is the ACK. The synthetic ethertypes are deliberate labels, not corruption	re-reading <code>synth_ack()</code> / <code>synth_resp()</code>	understated what the external evidence proves

#	mistake	how it was found	cost
32	Called defect 2 “refuted” when only three implementations were. All three tried to qualify the write at the producer, inside a stateful ALU that had no room; the constraint was real but the conclusion was too broad	asked what the write was actually for, and found it only had to let the next READ arm	the qualification moved to the consumer and cost nothing

The two that generalise:

A test that cannot fail proves nothing. Every checker in this directory has negative controls, and every one of them was mutation-checked — because four of the nineteen mistakes above were mistakes in the *checker*, not in the mechanism.

A silent success is not a success. Two of the four hardware traps compiled cleanly and produced no error at run time; one of them produced a *counter* saying the opposite of the truth. On this hardware, “it compiled and nothing complained” carries almost no information.

15 Summary

Defense 3 holds a DNP3 outstation’s TCP acknowledgement to a predetermined instant $t_{\text{ACK}} + D$, released independently of the response, using nothing but a crowd of recirculating dummy packets and the switch’s own strict-priority scheduler. It is implemented entirely in the data plane of an Intel Tofino-1, generates its own blocker tokens on chip, touches no packet byte, and requires no cooperation from the relay.

It does what it was built to do, on the one device tested. The CLRT — the cross-layer timing fingerprint that prior work uses to identify equipment — compresses from a 2.854 ms standard deviation to 0.012 ms, a factor of about 238. It does *not* become a constant: at $D = 16$ ms the 80 observations still take 18 distinct values with a $47 \mu\text{s}$ maximum. The campaign ran 480 physical transactions against a real SEL-751, **400 of them defended**, with zero fail-open events, zero queue drops and the acknowledgement-before-response ordering correct every single time.

And it is not a finished defense. It creates a new timing feature, READ→ACK, which is more separable from native traffic than the CLRT it removes — at *every* value of D , including the sub-threshold arm that collapses nothing. Equation (8) says why: while D is below the native CLRT the total is conserved and the defense merely redistributes it. At the D that conceals fully, the output does not resemble any physical device. Both facts were measured here, not argued, and both are stated in §12 rather than left for a reviewer to find.

And it is now correct as far as the switch can show. An external audit confirmed three defects (§7.5); all three are repaired and each is validated on silicon. The response marker is authorised before it writes (R1, validated across 1 920 live transactions and Gate 4 case F); the fail-open retire is generation-qualified through a second register (R2, the path now crediting all 64 tokens to the budget instead of 1); and a host-injected $0 \times 88\text{C}1$ frame is dropped before it enters the queue (R3, shown with an in-switch injector). A fail-open K -sweep pinned the second defect down: a single *native* budget-zero token clears the tag at $K = 1$, giving the $1/K - 1$ termination cascade, which R2 corrects to $K/0$ with the tag preserved — while the *cross*-transaction reach of that write still needs a generation-wrap coincidence and stays model-checked. **Every measurement in §10–§11 predates the repairs**, and the one thing still unshown is the repairs against a *real network attacker* rather than the switch’s own generated frames (§12).

The next experiment is the one comparison this report cannot make: an iso-latency Defense 2 arm at $G \approx D + \text{native CLRT}$, the only way to say whether holding the acknowledgement or holding the response is the better trade.