

Programmable In-Network Timing Obfuscation for DNP3

Anonymous Submission

Abstract—An adversary watching a power-grid control network can tell which device it is looking at from how long that device takes to answer, because vendors implement the same function differently. Hiding that timing is challenging. The relay cannot be modified, and existing traffic obfuscation was built for a different feature: it reshapes flow-level sizes and inter-packet times, whereas the feature that identifies the device here is one interval inside a single transaction. It also hides the packets it adds behind encryption, which these plaintext networks do not have.

To change the timing without touching either endpoint, we put the defense in the network. A programmable switch in front of the outstation holds the device’s acknowledgment and its response, and releases each at a deadline computed from the same instant, the moment the outstation acknowledges the request. Neither release depends on how long the device took, so the interval an observer measures is a value the operator sets. The switch forwards the packets it received, so both endpoints run unmodified and no DNP3 byte changes.

We implement this on an Intel Tofino-1 and evaluate it against one physical SEL-751A protective relay. The measured interval settles on the configured value: its median is the configured 4 ms, its interquartile range falls by more than two orders of magnitude, and its sample variance falls from 6.808 to 0.394 ms². A classifier trained on that interval before deployment is left at chance across the three transaction classes. An adversary who retrains recovers much of that from a second interval we do not target, and that residual bounds what the framework achieves. The evaluation covers transaction-class timing for a single device; we do not measure confusion between device models, and we do not claim the feature is removed for an adversary who accumulates many observations.

Index Terms—traffic obfuscation, device fingerprinting, DNP3, programmable switches.

I. INTRODUCTION

Device fingerprinting has been an essential step in cyber reconnaissance, allowing adversaries to reveal unique features of target networks and to design effective, stealthy attack strategies. These techniques are becoming increasingly critical in industrial control systems (ICSs) such as power grids, where adversaries often use IP-based control networks and computing devices within as entry points to inflict physical damage. Consequently, fingerprinting shifts the focus from visited websites and user biometric behavior to device models and types of control operations that are critical to ICS attacks.

In the 2015 attack that disrupted Ukrainian power grids and the Stuxnet attack that disrupted Iranian nuclear power facilities, it is widely believed that adversaries stay in their systems for at least 6 months to perform cyber reconnaissance [1], [2].

To disrupt device fingerprinting, many studies present network traffic obfuscation [3], [4]. Device fingerprinting targeting general computing environments generally relies on network-level features such as packet size and/or inter-packet latency observed from communication patterns [5]. Consequently, existing traffic obfuscation often focuses on (i) padding and splitting network packets, which hide or change the distribution of network packet sizes [3], and (ii) delaying network packets and adding dummy ones, which disrupt the inter-packet timing pattern [4], [6]. Since manipulating communication networks can introduce runtime overhead, recent works have begun to offload traffic obfuscation onto programmable network switches, which change communication patterns at much higher line rates than CPUs [7]–[9].

Unfortunately, these methods cannot be directly applied to ICS environments for two main reasons. First, device fingerprinting methods on ICS devices rely on different features from the ones in general computing environments. Instead of using network layer knowledge, they attempt to use behavior on physical devices to reveal device model or types of operations. For example, work in [10] uses the latency between the TCP acknowledge packets and the actual response from the same device to estimate execution time there. Since each vendors may choose different materials or algorithms to implement the same control functionality, this time can be accurate in fingerprinting those devices. Second, existing traffic obfuscation is performed over encrypted channels, which are necessary to mix the obfuscated and normal traffic. Unfortunately, ICS networks still rely on unencrypted traffic due to a wide range of legacy devices, despite ICS network protocols may provide security features. Directly padding or adding dummy packets can easily reveal obfuscated packets, exposing the the underlying traffic pattern to adversaries.

In this paper, we present a framework that schedules when selected packets of an ICS transaction leave the network, so that a timing feature an observer measures becomes a value we choose rather than a property of the device. A programmable switch at the outstation edge holds the acknowledgment and the response of a transaction and releases each at a deadline it computes when the transaction begins.

Because the switch forwards the same packets it received, the endpoints run unmodified firmware and every DNP3 byte reaches the master unchanged, which is what the plaintext

setting requires. The framework is general over which packets are held and what the released interval should be. We demonstrate it on two network-observable intervals: the cross-layer response time of a poll, which is the feature Formby et al. [10] fingerprint, and the master-visible timing of a select-before-operate control command. The second is our own case and is not their second feature, which is the physical operation time of the device rather than an interval read from the network. We evaluate the first on a physical relay. These are two case studies of one mechanism, not two systems.

This paper makes the following contributions:


- **Designs a timing-obfuscation framework** that replaces a selected within-transaction timing feature with a configured value by scheduling the release of real protocol packets. We derive the coupling between the acknowledgment hold, the response hold and the configured interval that any such schedule must satisfy (Section IV).
- **Realizes the framework on an Intel Tofino-1** using paired blocker and hold queues that sustain a hold until a deadline expires, with a read lane timed from the outstation's acknowledgment and a control lane timed from the request. Forwarding stays open when a response arrives after its scheduled release (Section V).
- **Evaluates the realization against a physical SEL-751A relay**, reporting how closely the released interval reaches its configured value, what the policy can and cannot cover, the latency the mechanism adds, the hardware resources it uses, and what an adaptive attacker still learns from the features that remain (Section VI).

II. BACKGROUND AND MOTIVATION

A defense that changes when a packet leaves has to be built on the events an observer can actually see. This section names those events for the two DNP3 transactions the paper covers. It then explains why the interval between two of them tells an observer something about the device, states which events we measured and which we did not, and describes what a programmable switch makes possible.

The DNP3 read and control transactions. DNP3 carries supervisory control and data acquisition traffic over TCP between a master and an outstation [11]. The master polls the outstation with a READ request, function code 1, and the outstation returns the requested measurements in an application-layer response, function code 0x81. A control command instead uses select-before-operate. The master sends a SELECT, function code 3, that names an output point, and the outstation returns a SELECT response. The master then sends an OPERATE, function code 4, which the outstation confirms with a solicited OPERATE response.

Because DNP3 runs over TCP, the outstation's transport layer acknowledges each request, and on the relay we evaluate it does so in a segment of its own, before the application answers. TCP does not require that: an acknowledgment may be carried on the response instead [12]. On this device the two are separate, and each transaction appears on the wire as a request, a transport acknowledgment carrying no DNP3



c: acknowledgment-to-response interval of the READ (CLRT).
O: master-visible OPERATE response-to-ACK interval.

Fig. 1. The two DNP3 transactions we measure, on the master-facing link. (a) A READ poll, whose cross-layer response time runs from the acknowledgment to the response. (b) A select-before-operate control, whose master-visible interval runs from the OPERATE's acknowledgment to its solicited response.

payload, and a DNP3 response, in that order, as Figure 1 shows. We use those three words for those three packets throughout the paper. The mechanism's eligibility condition is exactly this: it applies where the acknowledgment and the response are distinct packets, and forwards the exchange unchanged where they are not.

Why the interval between them is a fingerprint. The two answers come from different parts of the outstation. Its TCP stack emits the acknowledgment as soon as the request arrives, whereas its application emits the response only after the device has done the work the request asks for. Consequently, the interval between them, the cross-layer response time or CLRT, measures the outstation's own processing rather than the network path. It differs between device models and between the kinds of work a device is asked to do. For example, on our relay a control command takes 0.8 ms longer to answer than a poll, 2.937 ms against 2.116 ms, while the two requests are acknowledged 0.003 ms apart. Formby et al. used exactly this interval to tell device types and models apart in a substation network [10].

Our own measurements show the same effect on a single device: with no timing defense the three classes of transaction are already partly separable from timing alone, and Section VI reports the distributions.

The DNP3 function codes travel in plaintext and already name each operation, so the timing is not what reveals that an operation happened. What it adds is a second, device-derived signature of that operation.




Fig. 2. Observation model. The adversary sees only the master-facing link (shaded); the relay-facing link and the switch-internal lanes (dashed) are observed by neither the adversary nor us.

What we observe, and what we do not. Every quantity the paper measures is one of three packet timestamps taken on the master-facing link: the request, the acknowledgment and the response. The instants at which the switch releases a packet, and every event on the relay-facing link, were not captured. Where the paper names them, they come from the configuration and we say so, and no measurement in this paper observes physical breaker motion.

What a programmable switch makes possible. A programmable switch exposes a match-action pipeline, written in P4 [13], that parses headers, keeps state in registers, and places packets into output queues at line rate. The property that matters here is that such a pipeline can delay a packet and release it later without a host in the loop. Consequently, a switch placed between the master and the outstation can hold the acknowledgment and the response of a transaction and release them on a schedule an operator sets, changing neither endpoint nor any byte of the payload.

That possibility is what the rest of the paper develops. The next section states precisely which observer it is meant to defeat, and what we require of it.

III. THREAT MODEL AND OBJECTIVES

The adversary. We assume a passive observer on the master-facing link, the adversary of the fingerprinting attack we defend against [10]. Figure 2 shows where it sits. It records TCP and IP headers, segment sizes and arrival times. It does not modify or inject packets, compromise the master or the outstation, control the switch, or read the switch’s internal state. It does not see the relay-facing link, which in our deployment is inside the switch with no tap on it, and it does not see physical breaker motion.

Because DNP3 carries no transport security by default, we assume the traffic is unencrypted [14], so the observer also reads the DNP3 function codes directly. An adversary that probes, injects, or otherwise acts on the network is out of scope. The passive observer is the setting the traffic-obfuscation work we build on also targets.

What the adversary is looking for. Reconnaissance precedes the attacks that matter here. For example, false data

injection requires the attacker to know which devices it is dealing with [15]. Fingerprinting supplies that knowledge, and it does not need the payload, because the timing of the exchange carries it.

Following Formby et al. [10], two intervals are in scope. The first is the CLRT of a READ or of the SELECT phase of a control, which measures the outstation’s own processing as Section II describes. On our relay the first interval’s median differs by 0.9 ms between a poll and a control, against a request-to-acknowledgment interval that differs by 0.003 ms, so this is where the separation lives. The second is the master-visible interval between an OPERATE’s acknowledgment and its solicited response. The adversary also sees the request-to-acknowledgment interval, whose median is 0.56 ms on this device. The mechanism does not target it, and we report it as measured.

The second interval deserves a careful statement, because it is easy to claim more for it than we can support. Formby et al. estimate the time a device physically takes to complete an operation. However, they do so from a sequence-of-events timestamp carried inside the outstation’s application payload, and their alternative estimate from packet arrival times produced no usable result [10]. The interval we measure is a packet-timing feature of the solicited control response. It is not physical actuation time, and we do not measure physical actuation time anywhere. We also make no claim about the payload-timestamp feature, which a mechanism that moves packets without editing bytes could not affect in any case.

That narrower reading is still worth defending. On our relay the control lane answers 0.8 ms slower than the read lane, so the interval separates the two classes on its own, and Section VI-F shows an adversary recovering 0.651 balanced accuracy from exactly that kind of difference. RO2 therefore asks whether the interval becomes a policy value, not whether physical timing is concealed.

What our evaluation actually separates. Device fingerprinting is the application that motivates the work. What we evaluate is narrower, and we keep the two apart throughout. Our testbed has one outstation, so there is no second device to confuse it with and no device identification to measure. What we measure is whether the timing of an exchange still reveals *which class of transaction* it was, on that one device. That is a three-class problem over READ, SELECT and OPERATE, so chance is 0.333 balanced accuracy. A result here is therefore the suppression of a class-conditioned timing feature, and not a demonstration that two devices become indistinguishable.

Two adversaries, not one. We evaluate two conditions, which answer different questions, and we report both. The *fixed* adversary learns transaction classes from undefended traffic and applies that model unchanged once the defense is deployed. Because the profile is built before the defense is met, this is the adversary the reconnaissance setting describes.

The *adaptive* adversary collects obfuscated traffic and re-trains on it. It assumes more access than the reconnaissance setting describes, and it answers the question of what remains learnable once the defense is known: it recovers 0.651 bal-

anced accuracy where the fixed one is left at chance. Both conditions are ours, chosen to bracket the problem. The second is the one that limits our claim.

Research objectives. To make the goal checkable, we state it as five properties. The evaluation answers each in turn, and we reuse the labels throughout.

- **RO1 (read timing).** The visible CLRT of READ and of the SELECT phase of SBO is a configurable policy value rather than the outstation's own processing signature.
- **RO2 (control timing).** The master-visible OPERATE response-to-ACK interval stays at a policy value while the relay-facing hold is drawn from a configured codebook of 2, 6 and 12 ms.
- **RO3 (timing leakage).** The timing features carry less information about the transaction class, and a fixed adversary that trained on undefended traffic is reduced to approximately three-class chance, which for READ, SELECT and OPERATE is 0.333 balanced accuracy.
- **RO4 (release policy and coverage).** The release budget, the range of policies the mechanism can realize, the fraction of traffic it can cover, and the residual it leaves are characterized rather than assumed.
- **RO5 (cost and protocol preservation).** We measure the latency the mechanism costs and the frames and bytes it adds, and the external DNP3 workload still completes: every request is answered and every control operation returns a success status.

What is out of scope. Three boundaries follow from the model above and hold for every result we report. The framework does not hide the DNP3 function codes, which stay in plaintext, so RO3 concerns the timing channel and not the command semantics. It forwards every packet at its original size, which is what keeps it invisible in a plaintext setting. And the relay-facing link is not instrumented, so the per-transaction hold, drawn from a configured set of 2, 6 and 12 ms, and the release toward the outstation are configured rather than observed. That is why RO2 is stated in terms of the master-visible interval alone.

RO1 and RO2 both ask for an interval the operator chooses. Whether a switch can produce such an interval at all, and what it costs to do so, is a question about timing rather than about adversaries, and the next section answers it.

IV. DESIGN

The objectives of Section III require a selected interval to become a value the operator sets. We derive that from one transaction: a timing model that shows what has to change and why one arrangement of holds works where another does not, then the quantities that are free to choose and what bounds each of them, then the control transaction, which is timed from a different event. The hardware realization is deferred to Section V.

A. A timing model for one transaction

Consider one READ transaction, drawn in Figure 3. The master sends a request. The outstation answers it twice: first




Fig. 3. Release timeline of one READ transaction, drawn to scale but not to measurement: the time coordinates are illustrative, chosen so that an arrival and its emission are separable on the page. Both deadlines are computed from the acknowledgment's arrival t_A . In (a) the response arrives in time and the master observes the configured CLRT_{new}; in (b) it arrives late. The response hold $e_R - t_R$ is drawn in (a) because it is implied by the schedule, not configured. Marker shape identifies where an instant could be observed: filled circles lie on the master-facing link, which is the only place this work captured anything, and open squares are switch-side instants that were not captured.

its transport layer acknowledges the request, and later its application layer returns the response with the data. We write t_A for the instant the acknowledgment reaches the switch and t_R for the instant the response reaches it. The switch forwards each packet onward at some later instant, e_A and e_R respectively. The master therefore observes the interval $e_R - e_A$, up to the propagation and capture delay of the two packets on the link between them.

On our relay that interval has a median of 2.116 ms for a poll and a spread of 2.6 ms, and it is what carries information about the device:

$$\text{CLRT}_{\text{original}} = t_R - t_A, \quad (1)$$

the time the outstation itself spends between acknowledging the request and producing the answer. It is the feature Formby et al. use to tell device models apart [10], and Section II explains why; here it is simply the quantity we intend to remove.

To change what the master sees, the switch holds each packet. We write

$$D_A = e_A - t_A \quad (2)$$

for the *acknowledgment hold*, and we refer to $e_R - t_R$ as the *response hold* without giving it a symbol of its own, because the schedule below fixes it rather than letting us choose it. Substituting into the interval the master observes gives the identity that governs the whole design,

$$\text{CLRT}_{\text{new}} = e_R - e_A = \text{CLRT}_{\text{original}} + (e_R - t_R) - D_A. \quad (3)$$

Equation (3) makes the difference between two kinds of defense concrete. If the two holds are constants, their difference is a constant, and the observed interval is the original interval moved by that constant. The distribution slides along the axis and keeps its width, which on our relay means a spread of 2.6 ms wherever it is moved to. Consequently, an observer who has seen the device before can still recognize it by the shape. To remove the device's own contribution the holds cannot both be constants: the response hold must absorb the variation in $\text{CLRT}_{\text{original}}$.

To remove the device's contribution, we compute both releases from the same instant, t_A . When the acknowledgment arrives, the switch computes both release instants immediately,

$$e_A = t_A + D_A, \quad e_R = t_A + D_A + \text{CLRT}_{\text{new}}, \quad (4)$$

where CLRT_{new} is the interval we want the master to observe. We call this value the *configured* CLRT_{new} , and we call what the master actually records the *measured* CLRT_{new} ; Section VI reports how far apart the two are. Because both release instants derive from t_A and neither derives from t_R , the term $\text{CLRT}_{\text{original}}$ cancels in (3) and the observed interval is the configured one. Consequently, the response hold that results is not chosen but implied,

$$e_R - t_R = D_A + \text{CLRT}_{\text{new}} - \text{CLRT}_{\text{original}}. \quad (5)$$

Equation (5) is the reason the three quantities cannot be picked independently. Once the acknowledgment hold and the configured CLRT_{new} are fixed, the response hold follows from whatever the device happened to do on that transaction. Consequently, it varies from one transaction to the next exactly as much as the device does.

This holds only while the response exists in time to be released on schedule. The switch can delay a packet but cannot emit one it has not received, so (4) requires

$$t_R \leq t_A + D_A + \text{CLRT}_{\text{new}}. \quad (6)$$

Note that the acknowledgment may leave before the response arrives. That ordering is normal and does not by itself imply a lost packet. When (6) fails, the deadline is already in the past when the response appears and the switch forwards it immediately. The master then observes an interval larger than the configured value, as in Figure 3(b). One read-lane exchange in a thousand answers too late to be scheduled. We report those rather than excluding them, because they are the boundary of the mechanism and not a defect of the measurement.

The holds decide what the master observes between the two packets. They do not decide how long the master waits for

the answer itself, which is a separate quantity and the one an operational requirement would constrain. We write

$$D_R = m_R - t_R \quad (7)$$

for the *response latency*, where m_R is the instant the response reaches the master. Two terms make it up, $D_R = (e_R - t_R) + (m_R - e_R)$: the time the switch holds the response and the path from the switch to the master. The hold itself divides once more, into the wait to the release deadline, the interval ε from that deadline to the end of blocking, because the switch stops blocking a queued packet only once the queue gating it has drained, and the service that follows. Those three sum to the hold rather than adding to it. Consequently D_R is neither the response hold nor the configured CLRT_{new} . The outstation's own emission lies one hop before t_R and was not captured, so D_R omits that hop. Section VI-G reports what we measured of ε and on which build.

B. Choosing the holds and the configured interval

The model leaves three quantities to decide: the acknowledgment hold D_A , the configured CLRT_{new} , and, through (5), the response hold that follows from them. Each is limited by a different constraint.

The acknowledgment hold is limited in principle by the transport. Holding an acknowledgment delays the feedback the master's TCP uses to decide that a segment is lost. Consequently, a hold approaching the master's retransmission timeout, which RFC 6298 [16] floors at 1 s, would make the master resend a request the outstation had already answered.

That bound is less tight than it first appears, because the timer is adaptive. RFC 6298 [16] computes it from the round-trip times the stack has itself observed, so a hold applied uniformly is folded into those samples and raises the timeout rather than racing it; the standard notes the same effect from the attacker's side, as a way to inflate a sender's timer [16, §6]. What that argument does not cover is a hold that varies between transactions, which enters the variance term instead. Our policy holds by a constant, which is the case the adaptation absorbs.

In practice the transport bound is not the one that binds here. The mechanism sustains a hold with a finite resource described in Section V: a budget on blocker passes, not a wall clock. From that budget, an assumed service rate, a conservative bound on the outstation's own acknowledgment latency, measured detection and release terms and a safety margin, the control plane computes a horizon H of 30.8 ms and an admissible budget of 24.8 ms, and refuses to install a policy outside it.

That computation is an admission test in the control plane. It is not a deadline the data plane enforces, and the two must not be conflated. A pass budget bounds how many passes are *served*, and under strict priority a lower-priority reservoir can be starved while a higher-priority one is still being served, so one budget does not translate into one elapsed time for every queue. The sweep of Section VI-D shows the consequence directly: at a configured acknowledgment hold of 36 ms

the median request-to-acknowledgment interval saturates at 31.07 ms, well short of the configured value, while the median request-to-response interval for the same point is 40.58 ms and the median CLRT is 9.51 ms rather than the configured 4 ms. The acknowledgment reaches a ceiling; the response does not stop at it.

The saturation is therefore an observed transition for the acknowledgment lane at this configuration, and we report it as one. We do not claim it as a release deadline the hardware guarantees for every held packet, and Section VI-D reports the margin we observed. However, we do not claim the tested hold is universally safe: the master’s realized timer was never read back, and it is the one measurement that would settle the transport question directly.

The configured CLRT_{new} is limited by the operation, and it is limited through D_R rather than directly: what an operator can accept is a bound on how long the answer takes to arrive, not on the gap between the two packets that carry it. We looked for a communications requirement that fixes such a bound for a DNP3 poll or a select-before-operate control on a distribution relay, and we could not verify one. The only numeric deadline we could verify from primary documentation is the relay’s own select-to-operate timeout, whose documented default is 1.0 s [17], and that timeout governs how long a select stays armed rather than how quickly a response must return. Timing classes defined for other protocols govern those protocols’ own messages and do not transfer to a DNP3 poll merely because the numbers are convenient.

We therefore treat the setting as a deployment choice rather than a compliance question. We state the value we tested, 4 ms, and report the response latency the mechanism produces, about 23 ms more than the device alone, so an operator with a requirement of their own can check it against those numbers. The two constraints are different: the transport bounds D_A because the master’s retransmission timer runs on the acknowledgment, and the operation bounds D_R because the answer is what the operation waits for.

The acknowledgment hold and the configured CLRT_{new} together set what the mechanism costs. Their sum

$$D = D_A + \text{CLRT}_{\text{new}} \quad (8)$$

is the release budget: by (6) it decides which transactions the switch can place on the schedule at all. A larger budget covers more of the device’s own response-time distribution and costs more delay on every transaction it covers: at $D = 24$ ms it covers 99.9% of them for about 23 ms of added latency.

Note that D is the budget, not the delay added to each exchange. A transaction the outstation answers quickly is held longer than one it answers slowly, so the added delay is at most D rather than equal to it.

Because the budget and the observable are different quantities, an operator can fix the cost and still choose the feature. Holding D at 24 ms, we installed configured values from 1 ms to 22 ms and the master observed each one; the delay the mechanism added stayed the same across all of them.

Section VI-D reports that sweep and the point at which the realized hold stops tracking the requested one.

The framework and the policy it carries are separate. The model above says nothing about what CLRT_{new} should be. We chose a constant, because a constant removes the device’s variation entirely and is the simplest policy to state and to check. A schedule could instead draw the configured value from a distribution, or make it a function of the operation, and the same mechanism would carry it. We implement and evaluate the constant setting, and we do not claim evidence for the others.

C. The control transaction is timed from the request

A select-before-operate control command differs in one way that matters here: the framework holds the command itself before it reaches the relay, so that the moment the master issued the command is separated from the moment the relay acts on it. The outstation’s acknowledgment cannot be the reference instant here, because it does not exist until after the command has been released. Any deadline measured from it would carry the length of the command hold into the interval the master observes.

To keep the hold out of what the master sees, we time the control lane from the request. When the OPERATE arrives at T_0 , the switch releases it toward the relay after an internal delay J . It places the acknowledgment and the solicited response the master sees at configured offsets from T_0 ,

$$e_{\text{ack}} = T_0 + A, \quad e_{\text{or}} = T_0 + R. \quad (9)$$

Consequently, the interval the master observes is

$$O = e_{\text{or}} - e_{\text{ack}} = R - A, \quad (10)$$

in which J does not appear. Consequently, the relay-facing release time can vary while the master-facing interval does not, and an observer cannot subtract a known offset to recover it.

The control plane checks that A exceeds the largest admissible J plus the outstation’s own acknowledgment latency before it installs the values, since a deadline the switch cannot honor would fail open rather than protect. In our experiments D_A is 20 ms, CLRT_{new} is 4 ms, and $R - A$ is also 4 ms, so the two lanes present the same value to the observer.

D. What the design does not change

The mechanism replaces one interval per lane and leaves the rest of the exchange as it was. The DNP3 function codes remain in plaintext and the direction of every packet remains visible. Consequently, an observer can still tell a poll from a control command by reading the packet. The request-to-acknowledgment interval still contains the outstation’s own acknowledgment latency, displaced by the constant D_A but not reshaped by it, so whatever that interval reveals it continues to reveal.

The switch forwards the packets it received without padding them, injecting new ones, or editing any DNP3 field, which is what the plaintext setting requires and which also means the




Fig. 4. How one pipeline realizes the schedule. Ingress classifies each transaction and computes its release deadlines. Blocker packets drained ahead of the held packet sustain a hold until its deadline expires. Solid lines are master-facing and observed; dashed lines are not, as in Figure 2.

mechanism changes no size or ordering feature. Section VI-F reports what an attacker who is allowed to retrain on the obfuscated traffic still recovers from the features that remain.

V. IMPLEMENTATION

Section IV says when each packet should leave. A forwarding pipeline has no way to wait: a P4 program runs to completion per packet, can read a clock, but cannot sleep or schedule a callback. This section explains how the switch releases a packet at a chosen instant anyway.

A. Holding a packet with queues

What the switch does have is a traffic manager with several queues per port and a strict-priority scheduler, the primitive with which SP-PIFO approximates programmable scheduling on today’s hardware [18]. A scheduler is a device for deciding what does *not* leave. To hold a packet without a timer, we use that property.

A READ uses *four* queues on one internal port, as Figure 4 shows. Two carry the packets we delay, the acknowledgment and the response, and above each of those sits a *blocker queue* in strict priority carrying small packets the switch generates itself. While blockers keep arriving, the packet beneath them is never served, and each hold ends when its own blocker queue runs dry. The response is not simply queued behind the acknowledgment: it has its own held queue and its own blocker gate, which is what lets the two releases be scheduled independently. In descending priority the ladder is acknowledgment blockers, held acknowledgment, response blockers, held response. Blockers circulate on an internal loopback port and never reach a cable, so neither endpoint can observe them. Each carries a tag naming its transaction and lane, so a leftover blocker cannot delay a later transaction.

It would be natural to read this as a timer built from a count: if a blocker takes τ to drain, K blockers hold a packet for about $K\tau$. The program does not measure time that way. To measure time without a timer, the program stores an absolute release instant when the acknowledgment arrives, quantized to a 256 ns grid, and every one of the 64 recirculating blockers compares the current timestamp against it. While the difference is negative the blocker is regenerated; once it is not, the blocker is dropped instead. The deadline decides when the packet leaves. The blockers only have to keep the queue non-empty until then, so the reservoir is sized for continuity rather than for duration.

The queue does not empty at the deadline, however. Each blocker tests the deadline on its own pass, so the blockers still queued when it passes have to be served once more before the queue runs dry and the held packet is served. The release therefore trails the deadline by the time that draining takes. Each lane drains its own queue, the acknowledgment’s above the response’s in the priority order, and Section VI-G reports both intervals as unmeasured. This is what makes the coupling of Section IV-A implementable: the response’s release instant is computed at the acknowledgment, so the resulting hold absorbs whatever the outstation does without the switch ever measuring it.

B. One READ transaction

The request arrives, a table classifies it by function code, and the program forwards it at once, recording that a transaction is in flight and which acknowledgment to expect. When the acknowledgment arrives at t_A , a median of 0.56 ms after the request on our relay, the program writes both release instants from that one timestamp, which is what puts them on the same instant t_A that Section IV-A requires. It does so once from an unarmed sentinel so a duplicate acknowledgment cannot arm a second deadline. The acknowledgment enters the hold queue and the packet generator fills the blocker queue. A response that arrives before its deadline queues behind the acknowledgment, and the two leave in order at their scheduled instants. A response that arrives after its deadline has nothing left to wait for, so the program forwards it at once, as in Figure 3(b). When the response has left, the transaction state is cleared.

Read and control responses can be in flight together. To keep the two independent, we give each treatment its own loopback port, its own hold and blocker queue pair, and its own registers. Sharing one hold queue would let whichever waited longer delay the other and corrupt the policy the second was meant to present.

C. Failing open

Every path through the program ends in a forwarding decision. A packet outside the protected session, an unrecognized request, an acknowledgment with no armed transaction, and a response whose transaction has retired are all forwarded unchanged. Such an exchange loses its protection but not its delivery.

Every hold is bounded by a budget on blocker passes: once a reservoir has spent its budget the packet is released and its state cleared. The budget counts passes, not elapsed time, so the horizon H of Section IV-B, 30.8 ms for the configuration we evaluate, is what the control plane computes from that budget and an assumed service rate, and is the bound it admits a policy against. It is not a wall-clock deadline the data plane keeps. With the timing mode off the same program forwards everything immediately, which is how we collected the Timing OFF arm. The program holds and forwards packets and does nothing else to them, so every byte the master reassembles is the byte the relay produced.

D. Resources and testbed

The program is about 1,560 lines of P4₁₆ [13] for the Tofino Native Architecture, compiled with the Barefoot SDE 9.13. It occupies all twelve ingress match-action stages of one pipeline and six egress stages, with 112 tables. State is small: one-entry registers reached by stateful ALU actions, one access per packet per register, which is what the architecture permits. One 25 Gb/s port carries the master and one 1 Gb/s port the relay; two further ports are loopbacks carrying no cable. The first loopback carries the four-queue read ladder the model of Section IV needs. The second carries the control extension, a further pair of queues, blockers above a held OPERATE, in its own scheduling domain so that the OPERATE reservoir cannot be starved by the higher-priority read reservoirs sharing a server. The combined program therefore uses six queues across two loopback ports: four for the read path and two for the control path.

Every policy quantity is a runtime table entry. Consequently, an operator changes policy by writing entries rather than loading a program. The loaded build realizes the dual-deadline schedule only: its arming table admits that mode alone, so the acknowledgment-only and response-only limits of the model are properties of the design and not settings of this build. Configuring them left both packets unheld, and the measured interval was the outstation's own.

The control plane, about 890 lines of Python, brings up ports and queues, installs the packet-generator configuration, writes the timing parameters and reads them back. It also refuses configurations the data plane could not honor. It checks the admissible budget of Section IV-B and, on the control lane, that A exceeds the largest admissible J plus a bound on the outstation's acknowledgment latency. The outstation is an SEL-751A feeder protection relay, one of the DNP3 device classes whose timing Formby et al. fingerprint [10]; the master issues READ and select-before-operate operations to isolated control points through a guarded driver, so no control operation in the evaluation moves a breaker.

VI. EVALUATION

In this section, we evaluate the framework against the five objectives of Section III, RO1 to RO5, on the physical testbed of Section V.

A. What we measured, and where

Evaluation Cases. We compare two arms of the same program. In *Timing OFF*, the switch forwards the outstation's acknowledgment and response as they arrive; in *Obfuscated*, it holds them and releases them on the deadlines of Section IV. The two arms differ only in the release schedule. The switch is the only path between the master and the relay. Because all timestamps are taken on the master-facing link, every interval reported in this section is one a passive observer on the control network can record.

Concretely, the interval we report is the time between the arrival of the transport acknowledgment and the arrival of the application response, both observed on the master's own interface. The switch's internal instants, when the outstation's packets reach it and when it releases them, are on links that were not captured, so they are recovered from the program rather than measured.

Dataset. We report per DNP3 exchange, i.e., one request, the transport acknowledgment that follows it, and the application response.

The campaign holds 63,360 exchanges, half with the mechanism running and half without, collected over five hours in interleaved blocks so that drift over the window cannot separate the two arms. Every request was answered, and no exchange was retransmitted, malformed or excluded.

To characterize the release policy itself, we ran a separate sweep on the same testbed over 16 configured policies and three controls with the mechanism disabled.

B. Does the released interval reach its configured value? (RO1)

RO1. The read-path schedule of (4) does replace $\text{CLRT}_{\text{original}}$ with the configured CLRT_{new} . We compare the two arms over all 22 grouped runs as full empirical distributions. The shape of the Timing OFF distribution is what can make a single observation informative to an adversary.

Figure 5 shows the two READ distributions directly. Before obfuscation the interval is broad and multi-modal, spanning 1.01 to 83.46 ms with distinct peaks near 1.2, 2.1 and 4.0 ms. After it, 99.9% of transactions fall within 0.10 ms of the configured CLRT_{new} . The sample variance falls from 6.808 to 0.394 ms², a reduction of 94.2%, and the sample standard deviation from 2.609 to 0.628 ms. Figure 6 shows what the full-range view cannot. The released interval is not a single value but a narrow distribution about the configured value, and the Timing OFF arm places only 13.5% of its transactions anywhere in that window. Figure 7 repeats the comparison with no cutoff, so the one-sided tail that the main view collects into its overflow category remains visible in full.

In Figure 8(d), we show the median $\text{CLRT}_{\text{original}}$ is 2.116 ms for READ and 2.050 ms for SELECT, with interquartile ranges of 2.778 ms and 2.697 ms. Under the mechanism, both medians are 4.000 ms and the interquartile range of each falls to 0.006 ms, a reduction of more than two orders of magnitude in spread.




Fig. 5. Measured READ CLRT before and after obfuscation. Bars give the percentage of that arm’s own transactions in each bin, on common 1 ms bins anchored at 0 ms with a linear ordinate and shared limits. Finite bins are half-open, $[lo, hi)$; the hatched category at the right holds every transaction at or above 15 ms, which is 0.409% of (a) and 0.027% of (b). The dashed line in (b) is the configured $CLRT_{new}$; the measured mean lies 0.012 ms from it. Because the configured value falls exactly on a bin edge, the obfuscated mass divides between the 3 to 4 and 4 to 5 ms bins, which is a property of the grid and not of the measurement; Figure 6 resolves it into one narrow mode. In (a) $n = 26,400$, mean 2.800 ms, sd 2.609 ms, variance 6.808 ms^2 ; in (b) $n = 26,400$, mean 4.012 ms, sd 0.628 ms, variance 0.394 ms^2 .

Figure 8(a, b) shows why the median alone understates the change. The Timing OFF distributions rise in a small number of discrete steps, so a single observation can be informative about the class. The obfuscated distributions are a step at the scheduled release. This is because the release instant is computed from the outstation’s acknowledgment and a fixed offset, and thus, it no longer depends on how long the device took.

The change is a replacement, not a shift. The standard deviation falls from 2.609 to 0.628 ms for READ and from 2.267 to 0.024 ms for SELECT, variance ratios of 0.058 (95% CI [0.019, 0.101]) and 0.0001 (95% CI [1.4×10^{-5} , 3.4×10^{-4}]) under a session-level bootstrap.

A defense that merely shifted the interval would keep the Timing OFF spread, because translating a sample leaves its variance unchanged: the counterfactual of the Timing OFF distribution translated onto the protected median retains a standard deviation of 2.609 and 2.267 ms, four times and ninety-four times what we measure. That counterfactual is analytical, computed from the Timing OFF samples themselves. The loaded program has no shifting mode, so it is not a measured arm and we do not present it as one. The spread collapses because both releases are computed from the acknowledgment’s arrival, as Section IV sets out. A shift keeps




Fig. 6. The released interval, resolved. The same measurements as Figure 5 on 0.005 ms bins within 0.10 ms of the configured $CLRT_{new}$ (dashed). The ordinate is still the percentage of that arm’s entire sample, not of the window, so the bars are directly comparable with the other figures. The window holds 13.53% of (a) and 99.86% of (b).




Fig. 7. The same measurements over the entire measured range. Identical 1 ms bins anchored at 0 ms and no overflow category, so every tail out to the largest observation, 83.5 ms in (a) and 57.2 ms in (b), is drawn. This is the companion to Figure 5, which truncates at 15 ms to keep the body of the distributions legible.




Fig. 8. Measured interval per class. (a) READ and (b) the SELECT phase of SBO report the CLRT; (c) OPERATE reports the master-visible response-to-ACK interval, timed from the request. (d) The same measurements as quartiles. The abscissa is logarithmic and spans the full support.




Fig. 9. Per-run medians and interquartile ranges, in acquisition order. Panel (b) uses a magnified ordinate spanning 0.20 ms. The runs are one campaign, not independent replications.

the shape and moves it; the measurement concentrates instead.

Low spread at the wrong interval would not be normalization, so we also report where the distribution sits: both medians are the configured 4.000 ms, and the interquartile range of 0.006 ms is two orders of magnitude below the 1 ms step between neighbouring policy settings in the sweep of Section VI-D.

A tail remains, and it is one-sided: the largest obfuscated READ interval is 57.182 ms, and fewer than one READ exchange in a thousand departs from the scheduled release by more than 1 ms, while none falls materially below it. The asymmetry is the mechanism, not the data. A response that reaches the switch after its scheduled release cannot be moved backwards. Consequently, it is forwarded on arrival and lands above the configured value, and nothing can place one early. These are the coverage boundary of RO4 seen from the other side. The replacement therefore covers the exchanges the budget can schedule: the released interval is the policy value, and the transactions that miss it are counted rather than hidden.

Stability within the Campaign. In Figure 9, we plot each grouped run in acquisition order. The obfuscated median sits at the scheduled release in every run and for every class, while the Timing OFF medians stay separated. Consequently, the effect is not an artifact of one run or of drift over the five hours.

C. Does the control lane behave the same way? (RO2)

RO2. The control-path rule of (9) holds the master-visible response-to-ACK interval at the policy value. Because the

control path is timed from the request, its observable is $O = R - A$, in which the internal hold J does not appear, and the read-path budget D does not apply to it. The switch draws J per transaction from the configured codebook $\{2, 6, 12\}$ ms. As shown in Figure 8(c, d), the OPERATE median moves from 2.937 ms under Timing OFF to 4.000 ms under the mechanism, with the interquartile range falling from 2.827 ms to 0.006 ms, and the master-visible interval remained concentrated near the configured 4 ms policy value in every run. That is a result about the interval the master sees. Note that the relay-facing release at $T_0 + J$ was not observed (Section VI-G).

D. What range of settings does the mechanism support? (RO4)

RO4. The interval a passive observer records is a policy value the operator sets, over a usable range of settings. To measure that range on the switch itself, we installed 16 release policies in turn and ran a fixed DNP3 workload through each. Two series answer the question. The first holds the total budget at $D = D_A + \text{CLRT}_{\text{new}} = 24$ ms and moves the split between the two, which changes the visible CLRT_{new} while leaving the total unchanged. The second fixes the configured CLRT_{new} at 4 ms and raises D_A until the mechanism leaves its operating region. Each point is the median over its READ exchanges.

In Figure 10(b), we show that the measured CLRT_{new} follows the configured value along the identity line across the whole series: settings of 1, 2, 4, 8, 12, 16, 20 and 22 ms produce measured medians of 0.998, 1.999, 3.999, 8.001, 12.001, 16.000, 20.001 and 22.001 ms, while the end-to-end response time stays between 25.30 and 25.34 ms at every point. As a result, the leaking interval and the cost of the exchange can be set independently, which is what makes the mechanism a framework rather than one fixed setting. In Figure 10(c), the request-to-acknowledgment interval tracks D_A up to 30 ms, where the measured median is 30.571 ms and CLRT_{new} still sits at 4.001 ms.

Beyond that point, the acknowledgment interval saturates at about 31.07 ms and CLRT_{new} can no longer be held: it rises to 5.503, 7.498 and 9.507 ms at D_A of 32, 34 and 36 ms. The reason is that the reservoir that starves the queue carries a finite pass budget, so a hold ends when that budget is spent. The saturation sits close to the admission horizon H of Section IV, which is the agreement one would expect if the control plane’s assumed service rate is about right; it is not evidence that the data plane enforces H . From below, the range is closed by the outstation’s own response time, which the switch cannot anticipate: in Figure 10(a), the budget of 24 ms covers 99.900% of Timing OFF read-lane exchanges, and one in a thousand arrives too late to be held. The usable range is therefore 1 to 22 ms at a fixed budget, bounded below by the outstation’s own response time and above by the 31.07 ms saturation.

E. What does the mechanism cost? (RO5)

RO5. The mechanism’s cost is measured, and the external DNP3 workload still completes. As shown in Figure 10(d), the median response time rises by 22.7 ms for READ, 22.6 ms




Fig. 10. The release policy is programmable and bounded on both sides. (a) Timing OFF read-lane tail against the budget D and the horizon H . (b) At $D = 24$ ms the measured CLRT_{new} follows the configured CLRT_{new} while the response latency stays fixed. (c) Raising D_A saturates near H . (d) End-to-end response time per class.

for SELECT and 21.2 ms for OPERATE. Note that the added latency is close to, but below, the release budget, because it depends on how long the outstation itself took; it is a real operational cost that an operator would weigh against the polling period and any protocol timeout.

On the master-facing link each exchange carries 3.017 frames and 272.3 bytes, identical in the two arms, so the mechanism moves packets in time without adding a frame or a byte. Every request was answered with a well-formed response, and every control operation returned a success status. The measured cost is about 23 ms per exchange, with no added frame or byte on the master-facing link.

Timeout and Retransmission Headroom. Holding the acknowledgment consumes the master’s timers, so we measured what it consumed. On the master-facing captures we recovered every transport-level event: the request bytes were acknowledged in a separate segment every time, never piggybacked on the response, and no request was sent while earlier bytes were still unacknowledged. There was no retransmission, no reset and no duplicate acknowledgment in either arm.

The longest the master waited for its request to be acknowledged was 29.2 ms and the longest it waited for the response was 77.7 ms. We did not record the master’s kernel configuration, so its actual retransmission timeout is unknown; for scale, the minimum Linux documents is 200 ms and RFC 6298 recommends a 1 s floor, and no retransmission occurred. The receive timeout in our driver is 3 s, and it bounds one read rather than the whole transaction, so it is not a transaction deadline; we report the measured completion times against it without asserting that the two coincide, which would need application-level receive tracing we did not collect.

Select validity is the other timer a held control transaction could violate, since the master issues its OPERATE only 0.19 ms after the SELECT response reaches it and the mechanism delays that response: no obfuscated OPERATE exchange was refused for a stale select. All of this characterizes this operating point on this testbed, and would have to be rechecked for a master with a shorter timeout.

F. What does the attacker still learn? (RO3)

Setup. RO3 asks how much the timing features still tell an adversary about the transaction class. We treat this as a three-class problem over READ, SELECT and OPERATE, for which chance balanced accuracy is one third, with the two independent intervals, request to acknowledgment and acknowledgment to response, as features; the total response time is their sum and is excluded. We estimate mutual information for the CLRT in bits and compare it against a null built by shuffling class labels within each grouped run over 1,000 permutations, the reconnaissance setting of Formby et al. [10] being the one we defend against, which preserves the per-run class counts. Every fold leaves out one grouped run, and we report the spread of the 22 held-out-run scores as a range.

RO3. In Figure 11, we show what the mechanism does to the two measurable intervals before any classifier is applied. Under Timing OFF, the three classes are separated mainly

along the post-acknowledgment axis, with median intervals of 2.116, 2.050 and 2.937 ms for READ, SELECT and OPERATE against request-to-acknowledgment medians of 0.555, 0.555 and 0.558 ms, which barely separate the classes at all. Under the mechanism, that axis collapses: all three medians sit at 4.000 ms, and READ and SELECT overlap. What remains is a horizontal offset, with the request-to-acknowledgment median moving to 21.336, 21.239 and 20.650 ms. This is because the read lane is timed from the outstation’s acknowledgment and the control lane from the request.

In Figure 12, we show the two attacker models.

The fixed adversary trained on Timing OFF traffic reaches 0.651 balanced accuracy on held-out Timing OFF exchanges using the CLRT alone, and 0.733 using both intervals. Applied unchanged to obfuscated traffic, it falls to 0.333 and 0.334, which is chance. Mutual information between the CLRT and the class falls from 0.383 bits to 0.004 bits; the Timing OFF value lies far above its permutation null at the resolution of the test ($p = 0.001$ with 1,000 permutations), and the obfuscated value lies inside its null ($p = 0.096$). For the evaluated fixed Random-Forest attacker, then, three-class transaction identification falls to approximately chance, and the cross-layer response time stops carrying information about the class. Note that this holds for the two features we give it; the plaintext function code is outside a timing-only feature set and still names the operation.

The adaptive adversary, retrained on obfuscated traffic, remains at chance on the CLRT alone. However, it recovers to 0.651 when the acknowledgment latency is added. The reason is that the read and control paths are timed from different events, so their acknowledgment latencies differ. Therefore, an adversary willing to collect obfuscated traffic and retrain can exploit that difference; the confusion panels show its shape, with OPERATE separable while SELECT stays confused with READ.

We state the cost of that difference rather than leave it to be inferred. On the request-to-acknowledgment interval alone, balanced accuracy rises from 0.379 under Timing OFF to 0.662 under the mechanism: the interval carries almost nothing about the class before the mechanism runs, and is the most informative single feature afterwards. The mechanism therefore does not only fail to remove that feature, it creates it, because holding the acknowledgment for D_A displaces an interval that the two lanes compute from different events. Against an adversary that uses both intervals and is allowed to retrain, the net effect is a fall from 0.733 to 0.651, or 0.082 of balanced accuracy on a three-class problem whose chance is 0.333.

Both conditions are ours, and each answers a different question. Against the fixed adversary, which profiles the device from undefended traffic and applies that model unchanged, the released interval carries 0.004 bits and classification falls to chance. Against the adaptive adversary, which collects traffic through the mechanism and retrains, balanced accuracy is 0.651 and the mechanism’s net effect is the 0.082 above: the targeted feature is suppressed, the acknowledgment interval




Fig. 11. The two measurable intervals against each other, (a) Timing OFF and (b) Obfuscated. Under the mechanism the vertical interval of all three classes collapses onto the policy value, while OPERATE keeps a horizontal offset because the two lanes are timed from different events. Small marks are a subsample; the outlined marks are class medians over the full dataset.




Fig. 12. Transaction-class leakage under the two attacker models, leaving out one grouped run at a time. (a) Balanced accuracy; whiskers span the 22 held-out runs and are not a confidence interval. (b) Mutual information against a within-run permutation null. (c) and (d) Row-normalized confusion using both intervals.

is not, and the second largely offsets the first. Formby et al. [10] describe reconnaissance that profiles devices from undefended traffic, which is what motivates our first condition; their setting does not restrict an observer of this mechanism to an unchanged model, and we do not read it as doing so. Closing the residual requires both paths to be timed from the same event, which we leave to future work.

The targeted feature is therefore suppressed against the fixed adversary the threat model describes, and not against one that retrains.

G. Limitations

Each limitation below bounds one of the results above, and we name which.

One Device, One Campaign. The evidence is one SEL-751A behind one Tofino-1, in one approximately five-hour campaign. Consequently, every result from RO1 to RO5 is a statement about this device and this window. The runs are grouped collections rather than independent replications, so the spread across them is within-campaign variability.

What RO3 Establishes. The classification we evaluate separates three transaction classes on one outstation. It is not device-model identification, and with one outstation nothing here shows that two devices become indistinguishable. The result also characterizes the Random-Forest attacker we evaluated rather than every classifier, and the adaptive adversary of Section VI-F recovers 0.651 balanced accuracy from the difference between the two events, which is the bound on what RO3 achieves.

What RO2 Rests On. Only the master-facing link was captured, so the realized per-transaction hold J and the relay-facing release at $T_0 + J$ are configured rather than observed. RO2 is therefore a claim about the interval the master sees and reports no result per J .

What RO4 Observes. The admission horizon H is computed in the control plane rather than enforced by the data plane, so the saturation at 31.07 ms is an observed transition and not a proof that the two coincide. The sweep offsets likewise come from the configuration table rather than from a per-point readback.

The Residual Is Not Attributed. The switch’s own release instants were not captured in the campaign, and the program that produced it cannot supply them: the timestamp registers it declares are never written. What the measured distribution shows around the configured value is therefore a net difference between two releases, not a delay we can place inside the switch.

We measured the blocking interval separately, on an instrumented build. Adding two registers per lane to the same source, keyed on the outcome that marks a blocker terminating because its deadline passed, gives $\varepsilon = 1,705$ ns on the acknowledgment lane and 1,704 ns on the response lane, medians over twelve transactions, with a standard deviation of 9 ns and a range of 1,691 to 1,719 ns. The first blocker observes the expiry a median of 12 ns after the deadline, so detection is not a material term. The two lanes agreeing to within a

nanosecond is consistent with one reservoir mechanism gating both.

Three limits on that number. It comes from a *different build* than the campaign, so it characterizes the mechanism rather than explaining any particular exchange in the measured distributions. It ends at the last blocking action, not at the wire: ingress observes a packet arriving and not leaving, so the service after blocking ends is not included. And we did not test whether the instrumentation perturbs the timing it measures, which would need a campaign under the instrumented build compared against the frozen one. For scale, ε is about 4×10^{-4} of the 4 ms configured interval.

What RO5 Counts. The overhead of RO5 is measured on the master-facing link, where the switch’s internal blocker traffic never appears. The zero-retransmission result is bounded the same way: because the program suppresses a response retransmission that matches an already-seen transport position while its transaction is live, such a copy would be absorbed inside the switch and would not reach a master-facing capture.

That suppression raises a question we tested rather than argued: a response lost downstream of the switch is recovered by exactly such a retransmission, so a mechanism that discarded it would break loss recovery. Dropping the response inbound at the master, scoped to one connection, leaves it unacknowledged and the outstation retransmits on its own timer. All three retransmissions reached the master, at 3.21, 9.21 and 20.25 s. The mechanism forwards loss-recovery copies rather than discarding them, and the reason is a separation of timescales: a transaction is live for $D_A + \text{CLRT}_{\text{new}}$, about 24 ms here, while the outstation’s retransmission timer is about 3 s, so a copy arrives long after the transaction has retired. On an outstation whose timer approached the hold the two windows could overlap, and this result should not be carried to such a device without repeating the test.

The Tail Is Not Eliminated. A fixed read-path release budget leaves about one exchange in a thousand answering too late to be scheduled, and a comparable fraction of protected READ exchanges more than 1 ms from the configured value. RO1 is a claim about the distribution, not about every transaction in it.

VII. RELATED WORK

Device fingerprinting. Devices can be identified from what they emit, even without any access to the payload. Kohno et al. fingerprint a remote host from the skew of its clock as seen in TCP timestamps [19], and GTID fingerprints a device and its type from the timing signature of its packets [20]. These works define the threat we address; our objective, on the other hand, is to remove the timing features they rely on from the master-facing observation. Their target is the device type, whereas the classifier we evaluate separates three transaction classes on one outstation, so we suppress a feature their cross-layer response-time method uses without demonstrating that two devices become indistinguishable.

ICS and power-system fingerprinting. Formby et al. brought fingerprinting to control systems with the cross-layer

response time and the physical operation time, measured at the device, of DNP3 outstations [10], and Gu et al. added device physics as a feature [21]. Jeon et al. fingerprint SCADA devices passively without deep packet inspection [22]. The reconnaissance these techniques serve enables false-data injection and process-control attacks [15], [23].

Defenses in this setting have so far worked at the content and connectivity layer: DefRec virtualizes physical functions to present deceptive content [24], RAINCOAT randomizes data acquisition and injects decoy measurements [25], and a companion testbed evaluates such anti-reconnaissance approaches on grid infrastructure [26]. Compared to these, our framework works below the semantics they reshape and complements them: it changes when the wire carries the outstation’s answer, not what the answer says.

General traffic obfuscation. Traffic morphing transforms one class’s packet-size distribution into another’s [3], while Dyer et al. show that per-packet padding and fragmentation still leave exploitable features [4]. Tamaraw fixes packet size and rate at a bandwidth cost [6], and WTF-PAD pads inter-packet gaps adaptively [27]. These defenses are effective against the size features they target, and the attacks that follow them are correspondingly strong [5], [28].

The same size-and-timing channel identifies smart-home devices under encryption, where shaping is the countermeasure [29], and Pacer and NetShaper shape a tenant’s traffic to a schedule that bounds side-channel leakage in the cloud [30], [31]. Pacer and NetShaper each shape from a host or hypervisor that can add cover traffic, and the smart-home setting assumes an encrypted payload. Our setting differs in both respects rather than in the soundness of those methods: it carries a plaintext protocol whose leaking feature is one interval inside a request-response exchange, and no endpoint can be modified to help. Because the outstation cannot be changed, it adds no byte and runs entirely in the network.

Programmable data-plane traffic transformation. P4 defines the programmable parser and match-action model we use [13]. PIFO defines programmable scheduling at line rate [32], and SP-PIFO approximates it with strict-priority queues on today’s switches [18], which is the primitive our reservoirs use. Ditto pads and injects chaff to a fixed pattern at line rate [7], Minos morphs and schedules traffic on a programmable switch at low cost [8], and Securitas combines fragmentation and insertion, and realizes them across several data planes [9]. NetHide obfuscates topology in the data plane against link-flooding reconnaissance [33]. Closest to our setting, Hu et al. use programmable switches to enhance industrial-protocol security [34]. Our framework differs from these in the observable it addresses, i.e., the master-visible timing of a DNP3 transaction, and in holding packets rather than adding or reshaping them.

DNP3 timing and security. East et al. catalogue attacks on DNP3 and the absence of encryption that makes its semantics visible [14], and specification-based and state-based intrusion detection has been adapted to DNP3 [35], [36]. These do not change the timing an observer sees.

On the control path we report a different quantity from Formby et al.: the master-visible OPERATE response-to-acknowledgment interval, which is a protocol-response timing feature rather than the payload-timestamp estimate of physical operation time they use [10]. Section III states why the two are not interchangeable and what we therefore do not claim.

VIII. CONCLUSION

How long an outstation takes to answer is a signature an observer can profile, and the relay leaking it cannot be modified. This paper presents a timing-obfuscation framework that moves the problem into the network: a switch in front of the outstation holds the packets carrying the device’s timing and releases each at a deadline computed when the transaction begins. On the read path, which is READ and the SELECT phase of select-before-operate, both deadlines are computed from the same instant, the moment the outstation acknowledges the request; the OPERATE case is anchored to the request instead. Consequently, the interval an observer sees is the operator’s value rather than the device’s.

Evaluation on one physical relay shows the released interval concentrating on the configured value: its interquartile range falls from 2.78 to 0.006 ms and its standard deviation from 2.61 to 0.63 ms. A classifier trained on that interval before deployment drops from 0.651 balanced accuracy to three-class chance. The framework protects the exchanges its budget can schedule and forwards the rest untouched.

Two bounds limit the result. An adversary who retrains recovers to 0.651 by adding the request-to-acknowledgment interval, which the mechanism does not target, and the evidence is one relay, one campaign and a transaction-class evaluation rather than a comparison between device models. We also never captured the switch’s own release instants during the campaign, so the residual around the configured value is a difference between two releases rather than a delay we can place inside the switch; an instrumented build puts the post-deadline blocking interval at about 1.7 μ s, which bounds one of its terms but does not decompose it. In future work, we plan to evaluate the framework on outstations from several vendors and to capture the release instants directly on both switch links, which is what would attribute the residual rather than bound it.

ETHICS CONSIDERATIONS

This work required no human subjects and no personal data. The experiments ran on a laboratory testbed we own: one SEL-751A protective relay, one Intel Tofino-1 switch and one master host, on an isolated network with no connection to any operational system. No production grid, utility network or third-party device was contacted at any point.

The relay is physical equipment whose outputs can operate a breaker, so the control experiments were constrained accordingly. Every select-before-operate exchange in the evaluation addresses isolated control points through a guarded driver, and no control operation in this paper moved a breaker or energized a circuit. Every OPERATE in the evaluation is a

protocol exchange the relay accepted and acknowledged, not a physical actuation.

The mechanism we describe is defensive. It removes a timing feature that an adversary would use for reconnaissance, and it neither creates a new attack nor discloses a vulnerability in the relay, the switch or the DNP3 protocol. We report one effect that works against the defender, and we report it in full rather than omit it: holding the acknowledgment makes the request-to-acknowledgment interval more informative than it was, from 0.379 to 0.662 balanced accuracy for an adversary that retrains. An operator weighing deployment needs that number.

OPEN SCIENCE

The evidence and the analysis are reproducible from the raw captures. The packet captures, the extraction and statistics code, the figure generators, the pinned environment and the tests are archived together. One script rebuilds every number and the campaign figures from the captures; the histogram and model figures have their own generators and the manuscript has its own build, and the repository names all three rather than implying a single command. A hash manifest covers each published artefact, and a publication gate compares a fresh rebuild against what the paper contains, so a discrepancy between the manuscript and the evidence fails rather than passes silently. That gate is a byte comparison of rendered PDFs, which is stricter than numerical agreement: the numbers reproduce anywhere the pinned environment installs, while the last bits of a plotted coordinate need not agree across machines.

The switch program is the exact P4 source that ran, retained with its compilation record. We do not release the relay's firmware or vendor documentation, which are not ours to distribute. On acceptance we will publish the captures, the analysis pipeline and the switch source under an open license; the artifact is prepared for evaluation in the form described here.

REFERENCES

- [1] R. M. Lee, M. J. Assante, and T. Conway, "Analysis of the Cyber Attack on the Ukrainian Power Grid," Electricity Information Sharing and Analysis Center (E-ISAC) and SANS Institute, Defense Use Case, Mar. 2016.
- [2] R. Langner, "Stuxnet: Dissecting a Cyberwarfare Weapon," *IEEE Security & Privacy*, vol. 9, no. 3, pp. 49–51, 2011.
- [3] C. V. Wright, S. E. Coull, and F. Monrose, "Traffic Morphing: An Efficient Defense Against Statistical Traffic Analysis," in *Proceedings of the Network and Distributed System Security Symposium (NDSS)*, 2009.
- [4] K. P. Dyer, S. E. Coull, T. Ristenpart, and T. Shrimpton, "Peek-a-Boo, I Still See You: Why Efficient Traffic Analysis Countermeasures Fail," in *2012 IEEE Symposium on Security and Privacy (S&P)*, 2012.
- [5] P. Sirinam, M. Imani, M. Juarez, and M. Wright, "Deep Fingerprinting: Undermining Website Fingerprinting Defenses with Deep Learning," in *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security (CCS)*, 2018.
- [6] X. Cai, R. Nithyanand, T. Wang, R. Johnson, and I. Goldberg, "A Systematic Approach to Developing and Evaluating Website Fingerprinting Defenses," in *Proc. ACM SIGSAC Conf. on Computer and Communications Security (CCS)*, 2014, pp. 227–238.
- [7] R. Meier, V. Lenders, and L. Vanbever, "Ditto: WAN Traffic Obfuscation at Line Rate," in *Proceedings of the Network and Distributed System Security Symposium (NDSS)*, 2022.
- [8] Z. Wang, Q. Li, G. Xie, D. Zhao, K. Li, Z. Fan, L. Ma, and Y. Jiang, "Minos: A Lightweight and Dynamic Defense against Traffic Analysis in Programmable Data Planes," in *Proceedings of the 2025 USENIX Annual Technical Conference (ATC)*, 2025.
- [9] G. Xie, Q. Li, Z. Shi, G. Antichi, Y. Zhu, K. Li, C. Weng, S. Miano, Y. Jiang, and M. Xu, "Defending against Traffic Analysis Attacks with Flexible In-Network Obfuscation," in *Proceedings of the 23rd USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, 2026, pp. 2043–2063.
- [10] D. Formby, P. Srinivasan, A. Leonard, J. Rogers, and R. Beyah, "Who's in Control of Your Control System? Device Fingerprinting for Cyber-Physical Systems," in *Proceedings of the Network and Distributed System Security Symposium (NDSS)*, 2016.
- [11] IEEE, "IEEE Standard for Electric Power Systems Communications—Distributed Network Protocol (DNP3)," 2012.
- [12] W. Eddy, "Transmission Control Protocol (TCP)," 2022.
- [13] P. Bosshart, D. Daly, G. Gibb, M. Izzard, N. McKeown, J. Rexford, C. Schlesinger, D. Talayco, A. Vahdat, G. Varghese, and D. Walker, "P4: Programming Protocol-Independent Packet Processors," *ACM SIGCOMM Computer Communication Review*, vol. 44, no. 3, pp. 87–95, 2014.
- [14] S. East, J. Butts, M. Papa, and S. Shenoi, "A Taxonomy of Attacks on the DNP3 Protocol," in *Critical Infrastructure Protection III (IFIP AICT, Vol. 311)*. Springer, 2009, pp. 67–81.
- [15] Y. Liu, P. Ning, and M. K. Reiter, "False Data Injection Attacks against State Estimation in Electric Power Grids," in *Proc. ACM SIGSAC Conf. on Computer and Communications Security (CCS)*, 2009, pp. 21–32.
- [16] V. Paxson, M. Allman, J. Chu, and M. Sargent, "Computing TCP's Retransmission Timer," 2011.
- [17] Schweitzer Engineering Laboratories, *SEL-751A Feeder Protection Relay Instruction Manual*, 2013, pM751A-01-NB, 20130329 printing; settings table, `STIME01 select/operate time-out`.
- [18] A. G. Alcoz, A. Dietmüller, and L. Vanbever, "SP-PIFO: Approximating Push-In First-Out Behaviors using Strict-Priority Queues," in *Proc. 17th USENIX Symp. on Networked Systems Design and Implementation (NSDI)*, 2020, pp. 59–76.
- [19] T. Kohno, A. Broido, and K. C. Claffy, "Remote Physical Device Fingerprinting," *IEEE Transactions on Dependable and Secure Computing*, vol. 2, no. 2, pp. 93–108, 2005.
- [20] S. V. Radhakrishnan, A. S. Uluagac, and R. Beyah, "GTID: A Technique for Physical Device and Device Type Fingerprinting," *IEEE Transactions on Dependable and Secure Computing*, vol. 12, no. 5, pp. 519–532, 2015.
- [21] Q. Gu, D. Formby, S. Ji, H. Cam, and R. Beyah, "Fingerprinting for Cyber-Physical System Security: Device Physics Matters Too," *IEEE Security & Privacy*, vol. 16, no. 5, pp. 49–59, Sep. 2018.
- [22] S. Jeon, J.-H. Yun, S. Choi, and W.-N. Kim, "Passive Fingerprinting of SCADA in Critical Infrastructure Network without Deep Packet Inspection," arXiv:1608.07679 [cs.CR], Aug. 2016.
- [23] A. A. Cárdenas, S. Amin, Z.-S. Lin, Y.-L. Huang, C.-Y. Huang, and S. Sastry, "Attacks Against Process Control Systems: Risk Assessment, Detection, and Response," in *Proceedings of the 6th ACM Symposium on Information, Computer and Communications Security (ASIACCS)*, 2011, pp. 355–366.
- [24] H. Lin, J. Zhuang, Y.-C. Hu, and H. Zhou, "DefRec: Establishing Physical Function Virtualization to Disrupt Reconnaissance of Power Grids' Cyber-Physical Infrastructures," in *Proc. Network and Distributed System Security Symposium (NDSS)*, 2020.
- [25] H. Lin, Z. T. Kalbarczyk, and R. K. Iyer, "RAINCOAT: Randomization of Network Communication in Power Grid Cyber Infrastructure to Mislead Attackers," *IEEE Transactions on Smart Grid*, vol. 10, no. 5, pp. 4893–4906, 2019.
- [26] H. Lin, B. Shrestha, and Y.-C. Hu, "Cyber-Physical Testbed: Case Study to Evaluate Anti-Reconnaissance Approaches on Power Grids' Cyber-Physical Infrastructures," in *Proc. Workshop on Learning from Authoritative Security Experiment Results (LASER)*, 2020.
- [27] M. Juarez, M. Imani, M. Perry, C. Diaz, and M. Wright, "Toward an Efficient Website Fingerprinting Defense," in *Computer Security – ESORICS 2016*, 2016.
- [28] A. Panchenko, F. Lanze, J. Pennekamp, T. Engel, A. Zinnen, M. Henze, and K. Wehrle, "Website Fingerprinting at Internet Scale," in *Proc. Network and Distributed System Security Symposium (NDSS)*, 2016.
- [29] N. Aphorpe, D. Y. Huang, D. Reisman, A. Narayanan, and N. Feamster, "Keeping the Smart Home Private with Smart(er) IoT Traffic Shaping," *Proceedings on Privacy Enhancing Technologies (PoPETs)*, 2019.
- [30] A. Mehta, M. Alzayat, R. D. Viti, B. B. Brandenburg, P. Druschel, and D. Garg, "Pacer: Comprehensive Network Side-Channel Mitigation in the Cloud," in *Proceedings of the 31st USENIX Security Symposium*, 2022.
- [31] A. Sabzi, R. Vora, S. Goswami, M. Seltzer, M. Lécuyer, and A. Mehta, "NetShaper: A Differentially Private Network Side-Channel Mitigation System," in *Proceedings of the 33rd USENIX Security Symposium*, 2024.
- [32] A. Sivaraman, S. Subramanian, M. Alizadeh, S. Chole, S.-T. Chuang, Agrawal, Anurag, H. Balakrishnan, T. Edsall, S. Katti, and N. McKeown, "Programmable Packet Scheduling at Line Rate," in *Proc. ACM SIGCOMM Conference*, 2016, pp. 44–57.
- [33] R. Meier, P. Tsankov, V. Lenders, L. Vanbever, and M. Vechev, "NetHide: Secure and Practical Network Topology Obfuscation," in *Proc. 27th USENIX Security Symposium*, 2018, pp. 693–709.
- [34] Z. Hu, H. Lin, L. Waind, Y. Qu, G. Chen, and D. Jin, "Industrial Network Protocol Security Enhancement Using Programmable Switches," in *Proceedings of the IEEE International Conference on Communications, Control, and Computing Technologies for Smart Grids (SmartGridComm)*, 2023.
- [35] H. Lin, A. Slagell, C. D. Martino, Z. Kalbarczyk, and R. K. Iyer, "Adapting Bro into SCADA: Building a Specification-Based Intrusion Detection System for the DNP3 Protocol," in *Proceedings of the Eighth Annual Cyber Security and Information Intelligence Research Workshop (CSIIRW)*, 2013.
- [36] I. N. Fovino, A. Carcano, T. D. L. Murel, A. Trombetta, and M. Masera, "Modbus/DNP3 State-Based Intrusion Detection System," in *2010 24th IEEE International Conference on Advanced Information Networking and Applications (AINA)*, 2010, pp. 729–736.